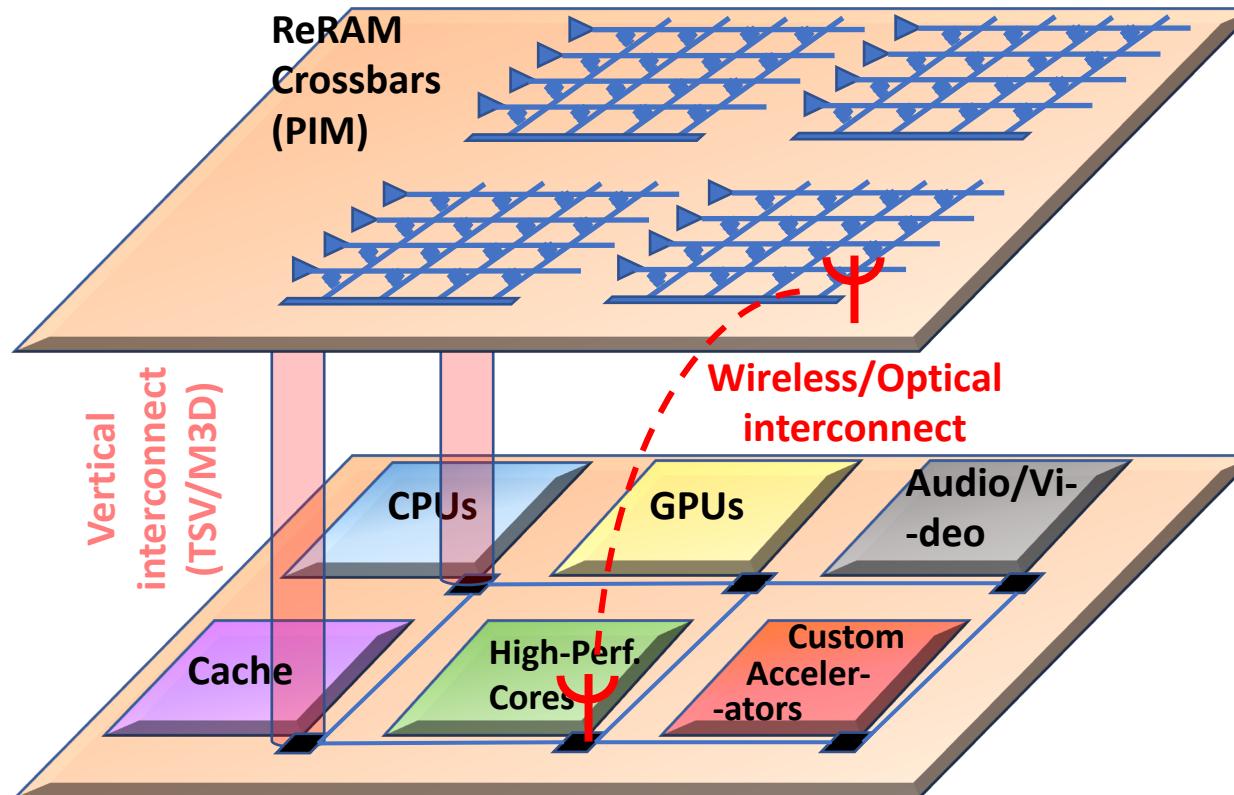
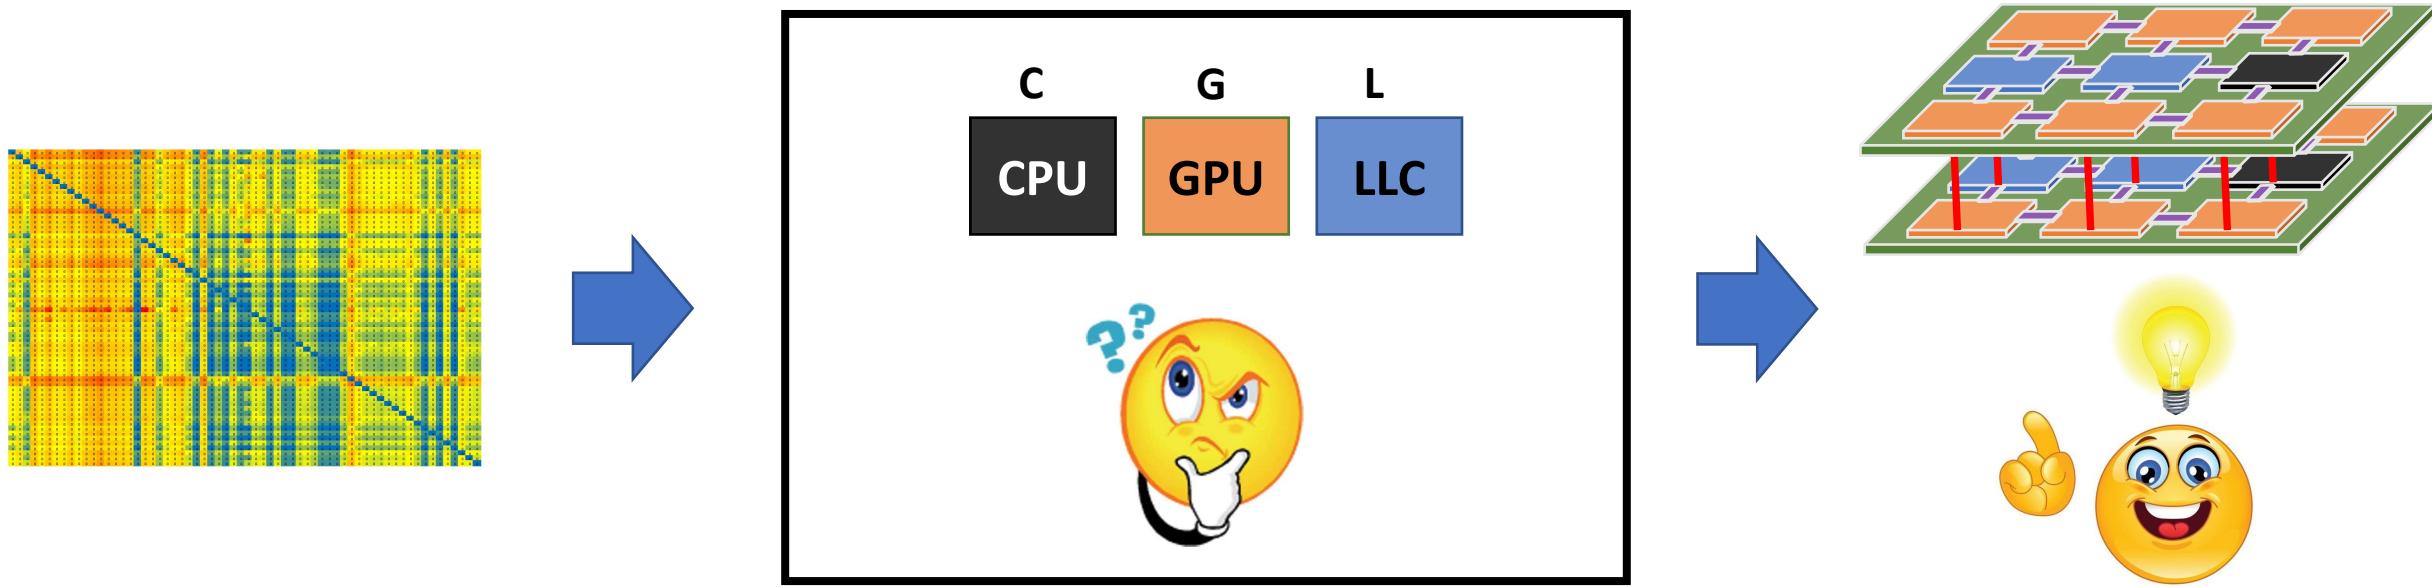


# Heterogeneous manycore design

- Move CPUs/GPUs/Accelerators on to the chip
  - Avoid off-chip data transfers
- High Performance, Low Power, More flexibility
- Examples: Apple 'A' series, Qualcomm Snapdragon



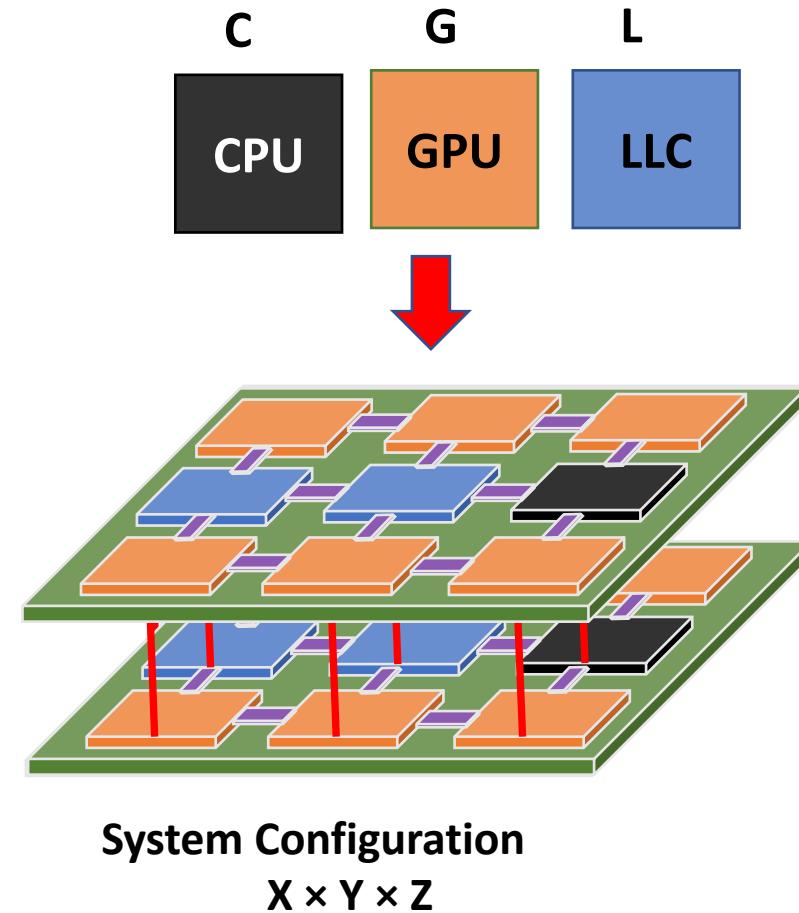
# Application-specific heterogeneous 3D manycore design



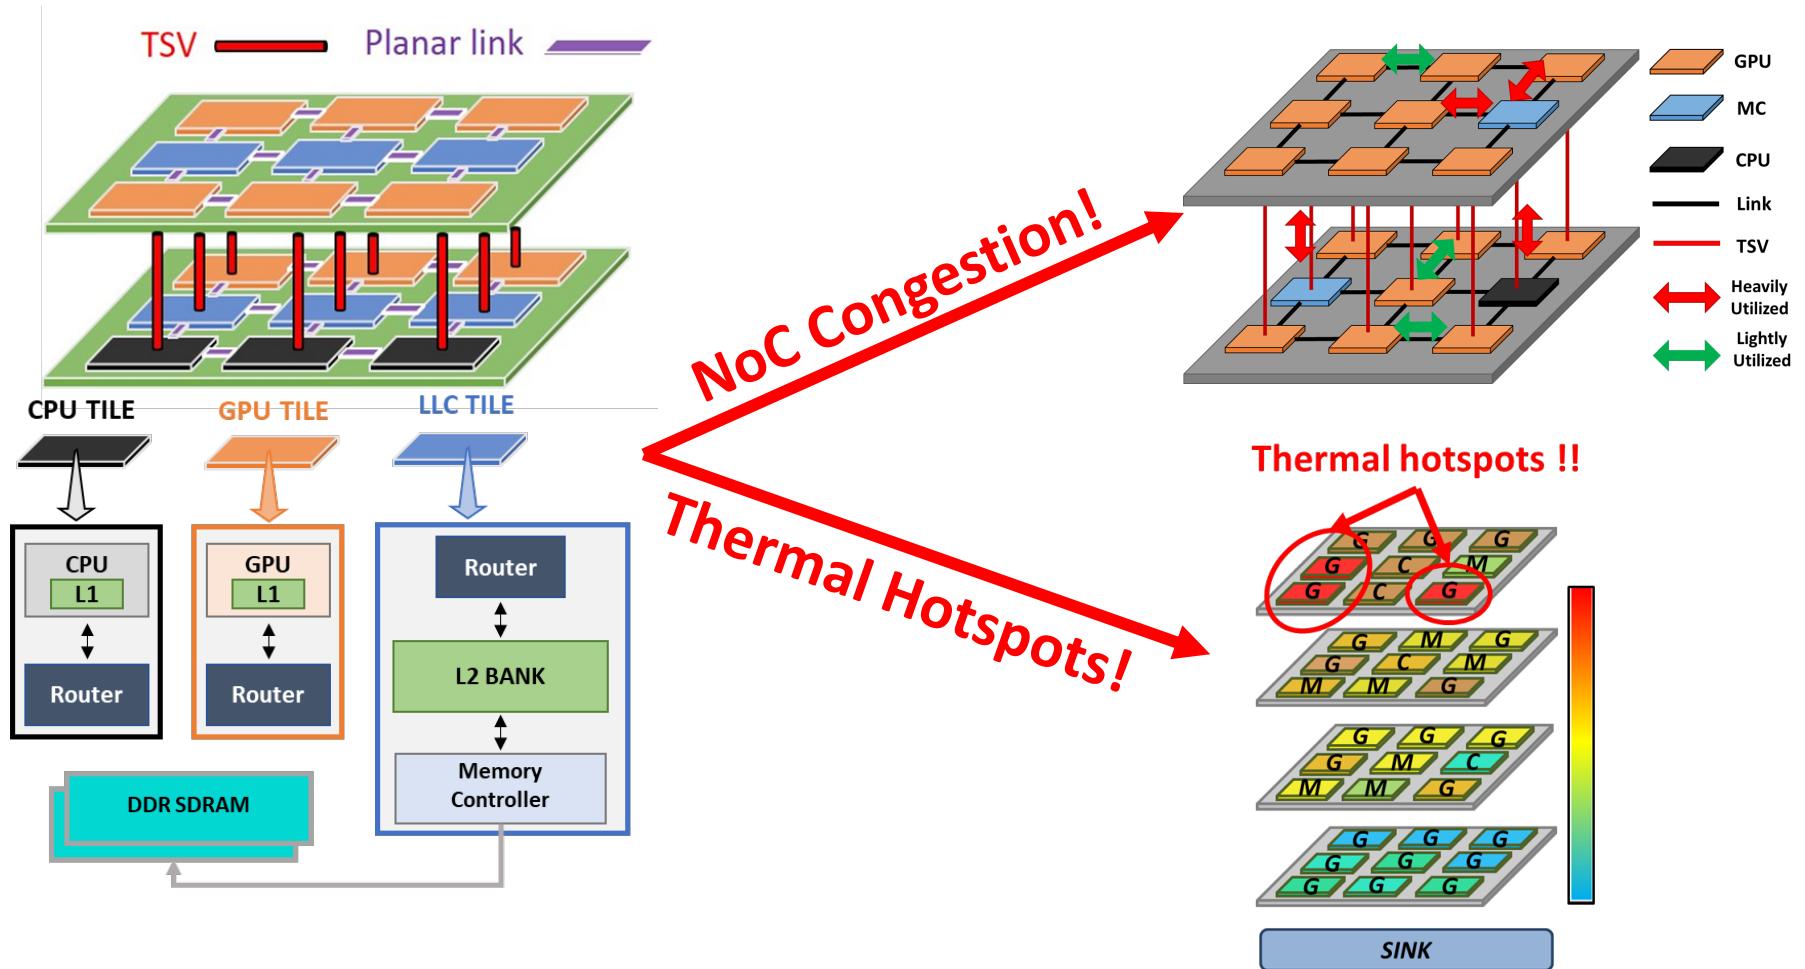
- Given an application behavior, what is the best architecture that we can design?
- Application behavior represented using traffic patterns/Power numbers
- Floor-planning problem

# 3D Manycore design

- Given:
  - $C$  CPUs,  $G$  GPUs,  $L$  LLCs and *mesh NoC*
  - System configuration ( $X \times Y \times Z$ )
- Goal:
  - Place CPU, GPU, LLCs and links
  - Satisfy different objectives simultaneously
    - *GPU throughput*
    - *CPU latency*
    - *Thermal, etc.*



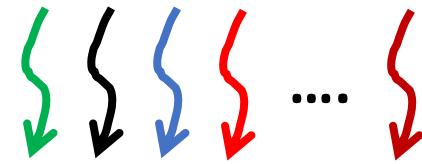
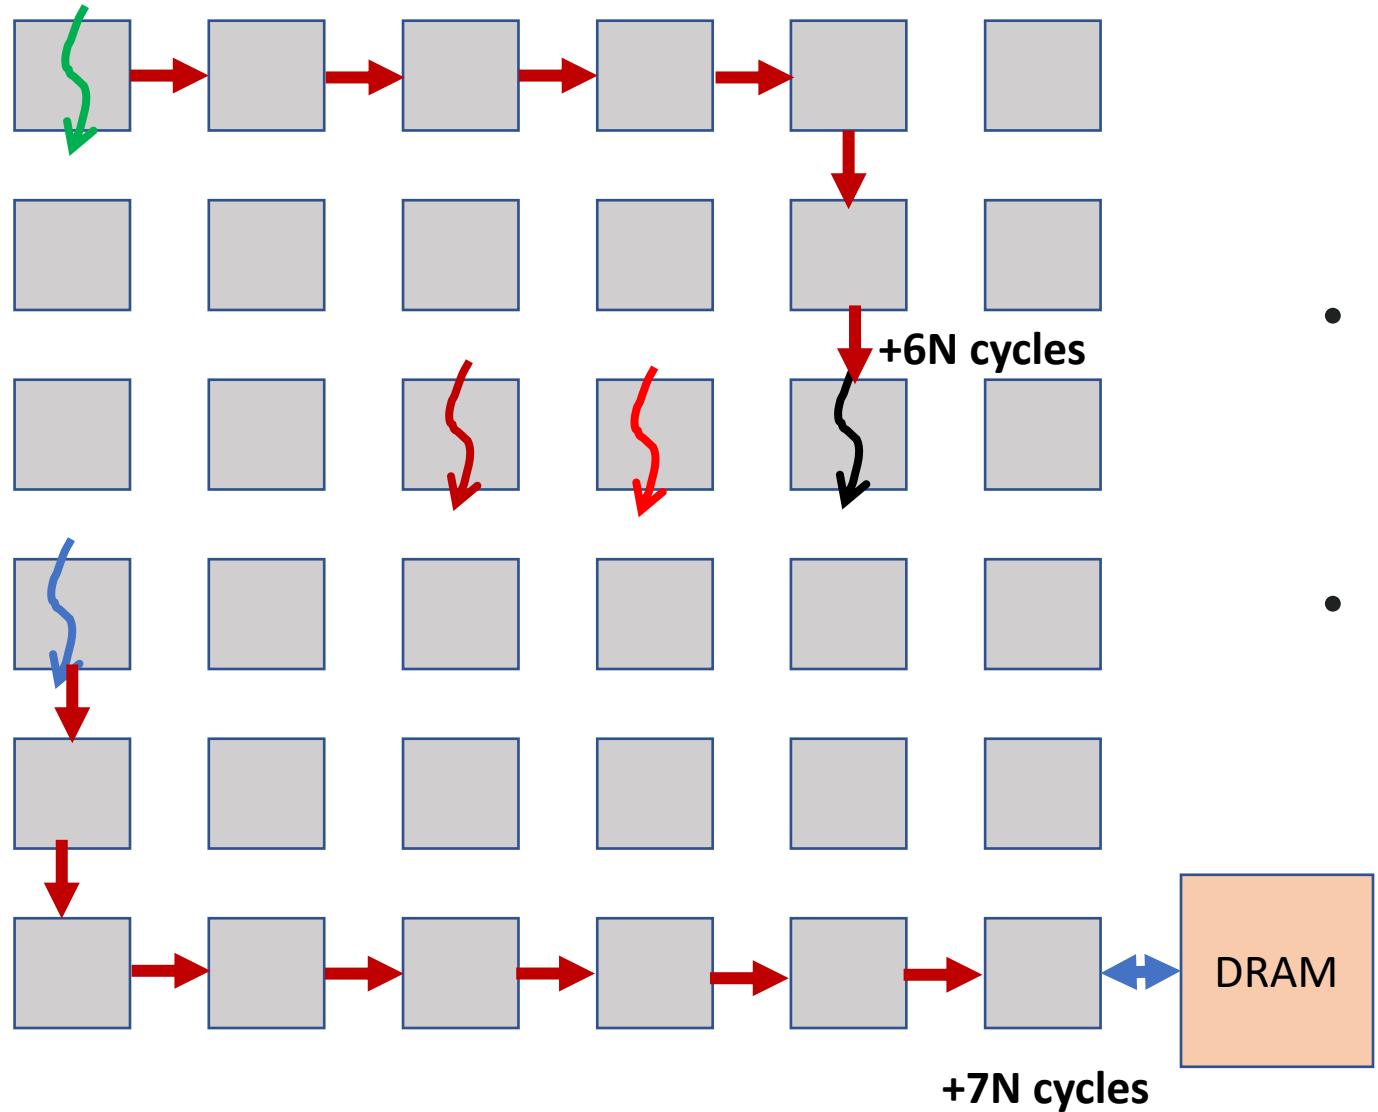
# 3D Manycore Design Challenges



- 3D manycore are a promising candidate

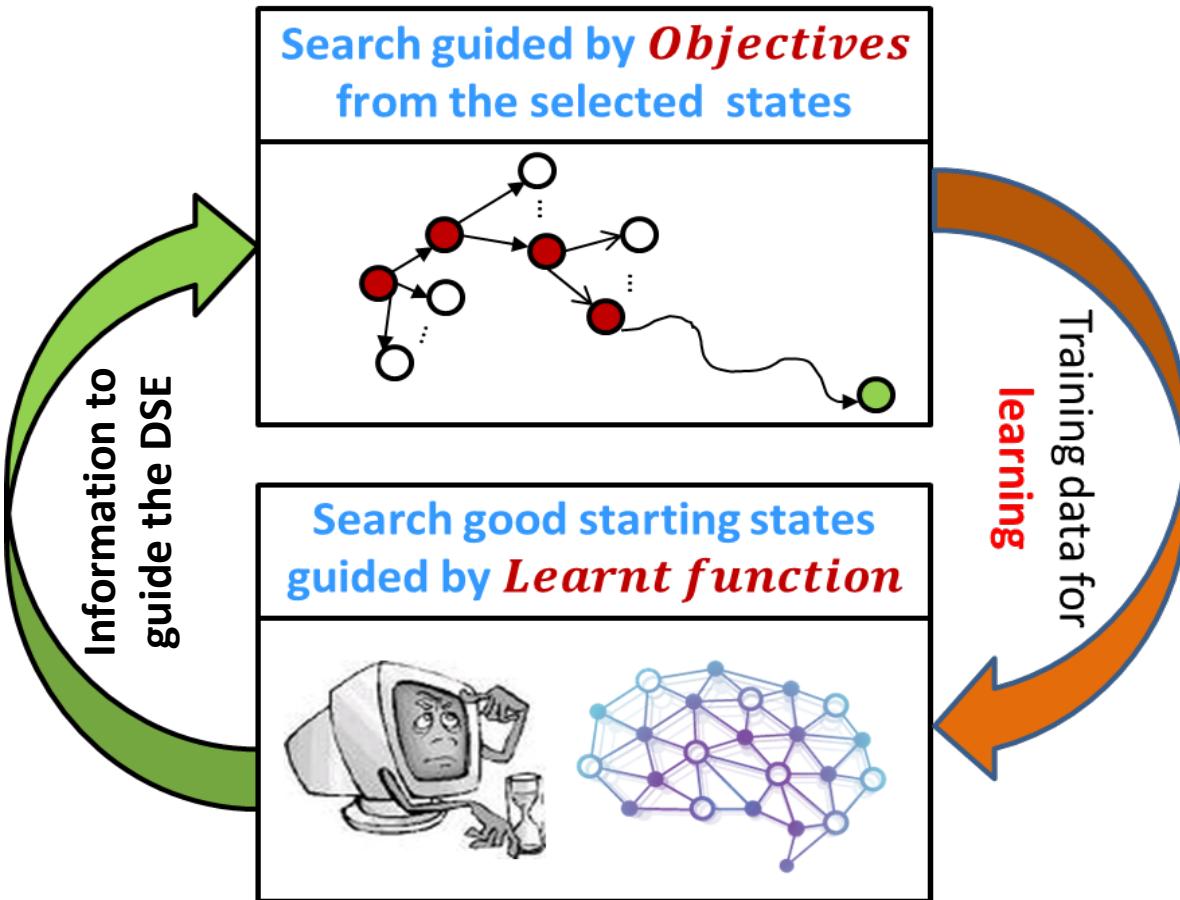
[1] B. K. Joardar, R. G. Kim, J. R. Doppa, P. P. Pande, D. Marculescu and R. Marculescu, "Learning-Based Application-Agnostic 3D NoC Design for Heterogeneous Manycore Systems," in IEEE Transactions on Computers, vol. 68, no. 6, pp. 852-866, 1 June 2019

# Similarity with task mapping



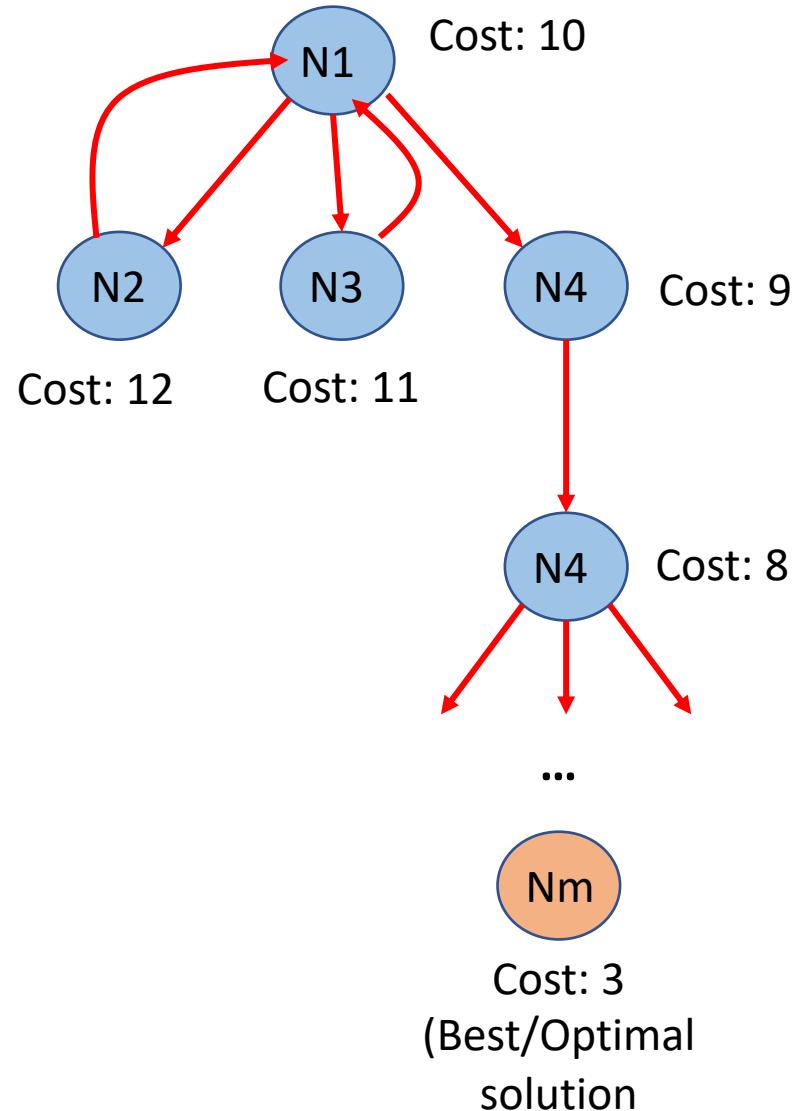
- **N tasks → N cores**
  - **Communication**
    - Performance/Execution time
  - **Thermal**
- **N! (N factorial) cases**
  - **Impossible to explore all possible solutions**
  - **How to map these tasks?**

# ML based DSE



- Use ML to guide the DSE
- More scalable
- One instance: Use ML to choose better starting points
  - MOO-STAGE [1]
  - Replace “Random restart” by “Guided Restart”

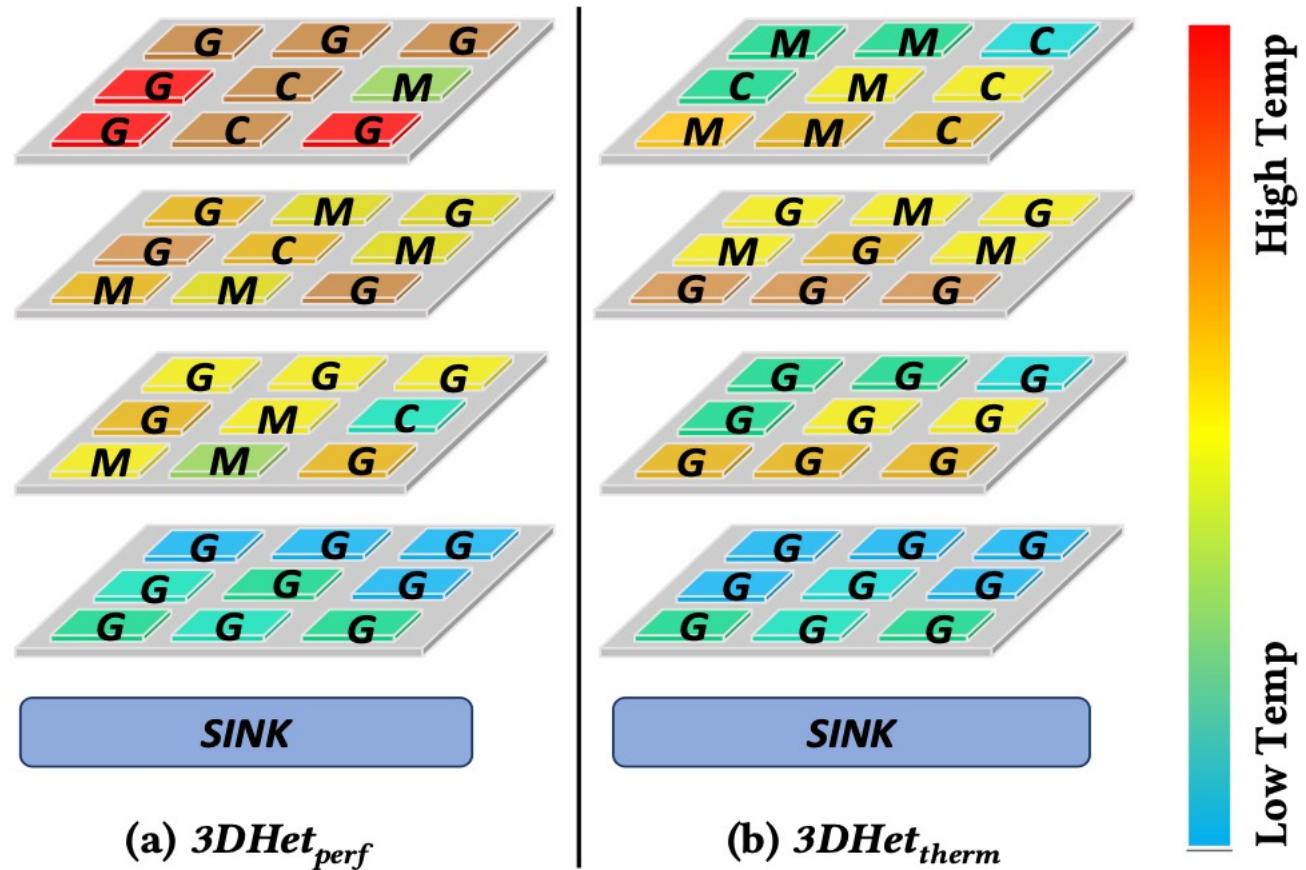
# Manycore Optimization



- **Similar to MOO-STAGE based NoC design problem**
- **Consider “Thermal only” and “Performance-thermal joint” optimization**

# Effect of core position on temperature

- By suitably mapping tasks to cores, we can reduce temperature
- High-power consuming tasks should be mapped near the sink
- Up to 15C cooler

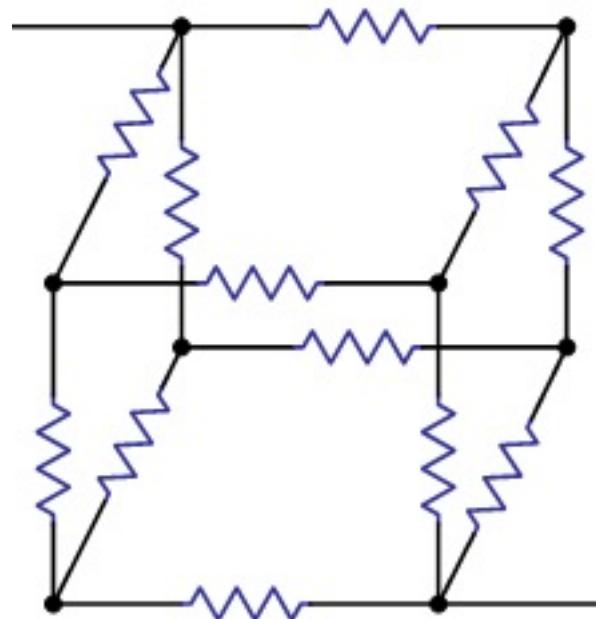


# Thermal-aware DSE

- **3D-specific objective: Temperature**

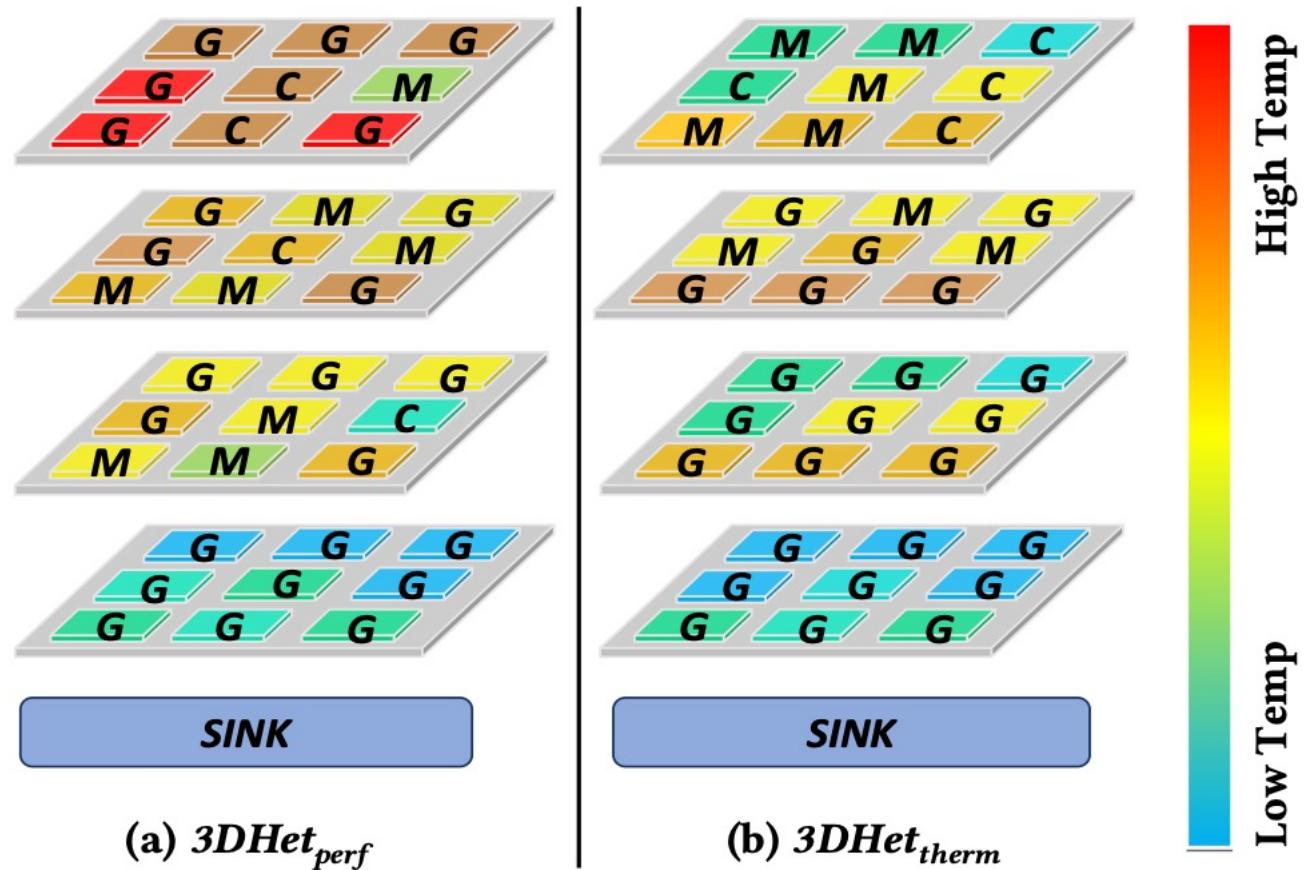
- Max. on-chip temperature

- $$T = \max_{n,k} \left\{ \sum_{i=1}^k \left( P_{n,i}(t) \sum_{j=1}^i R_j \right) + R_b \sum_{i=1}^k P_{n,i}(t) \right\} * T_H$$

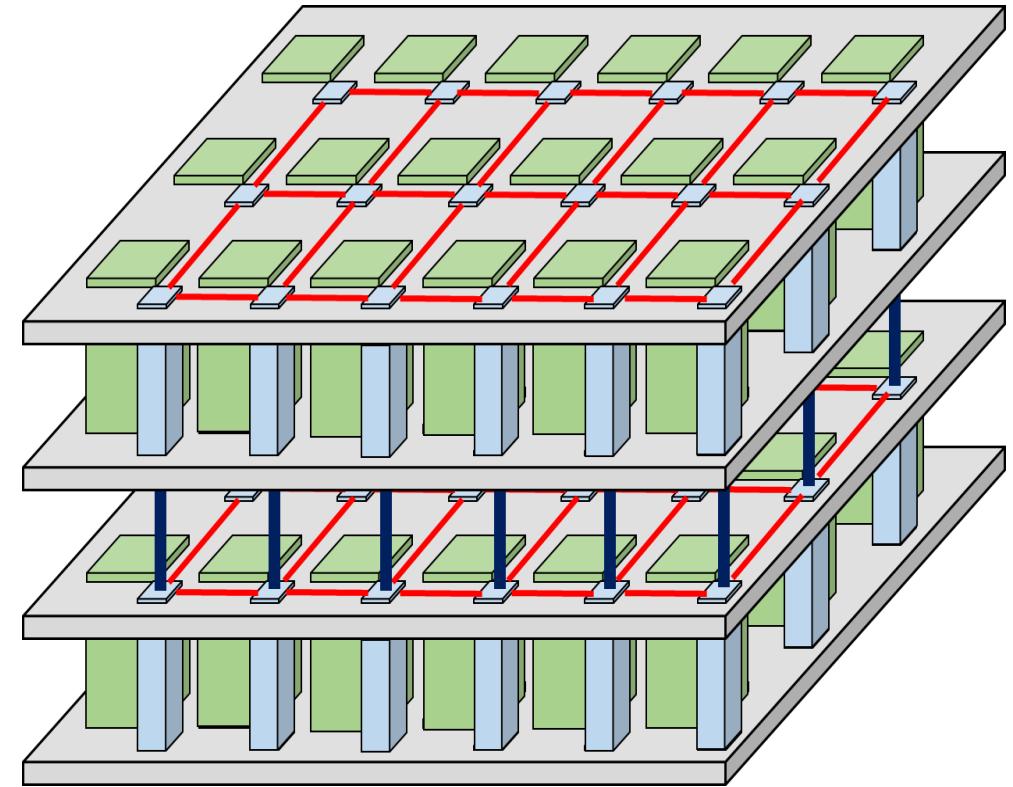
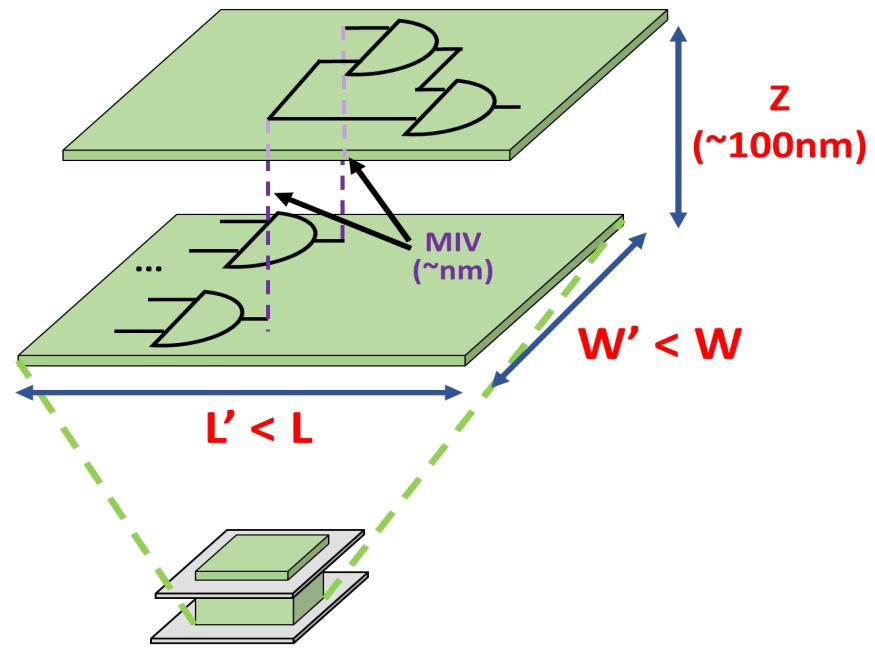


# Thermal-aware manycore design with TSV

- 3D architecture by varying core placement
- Reduce temperature
  - Place higher power cores near the sink
- Mesh NoC

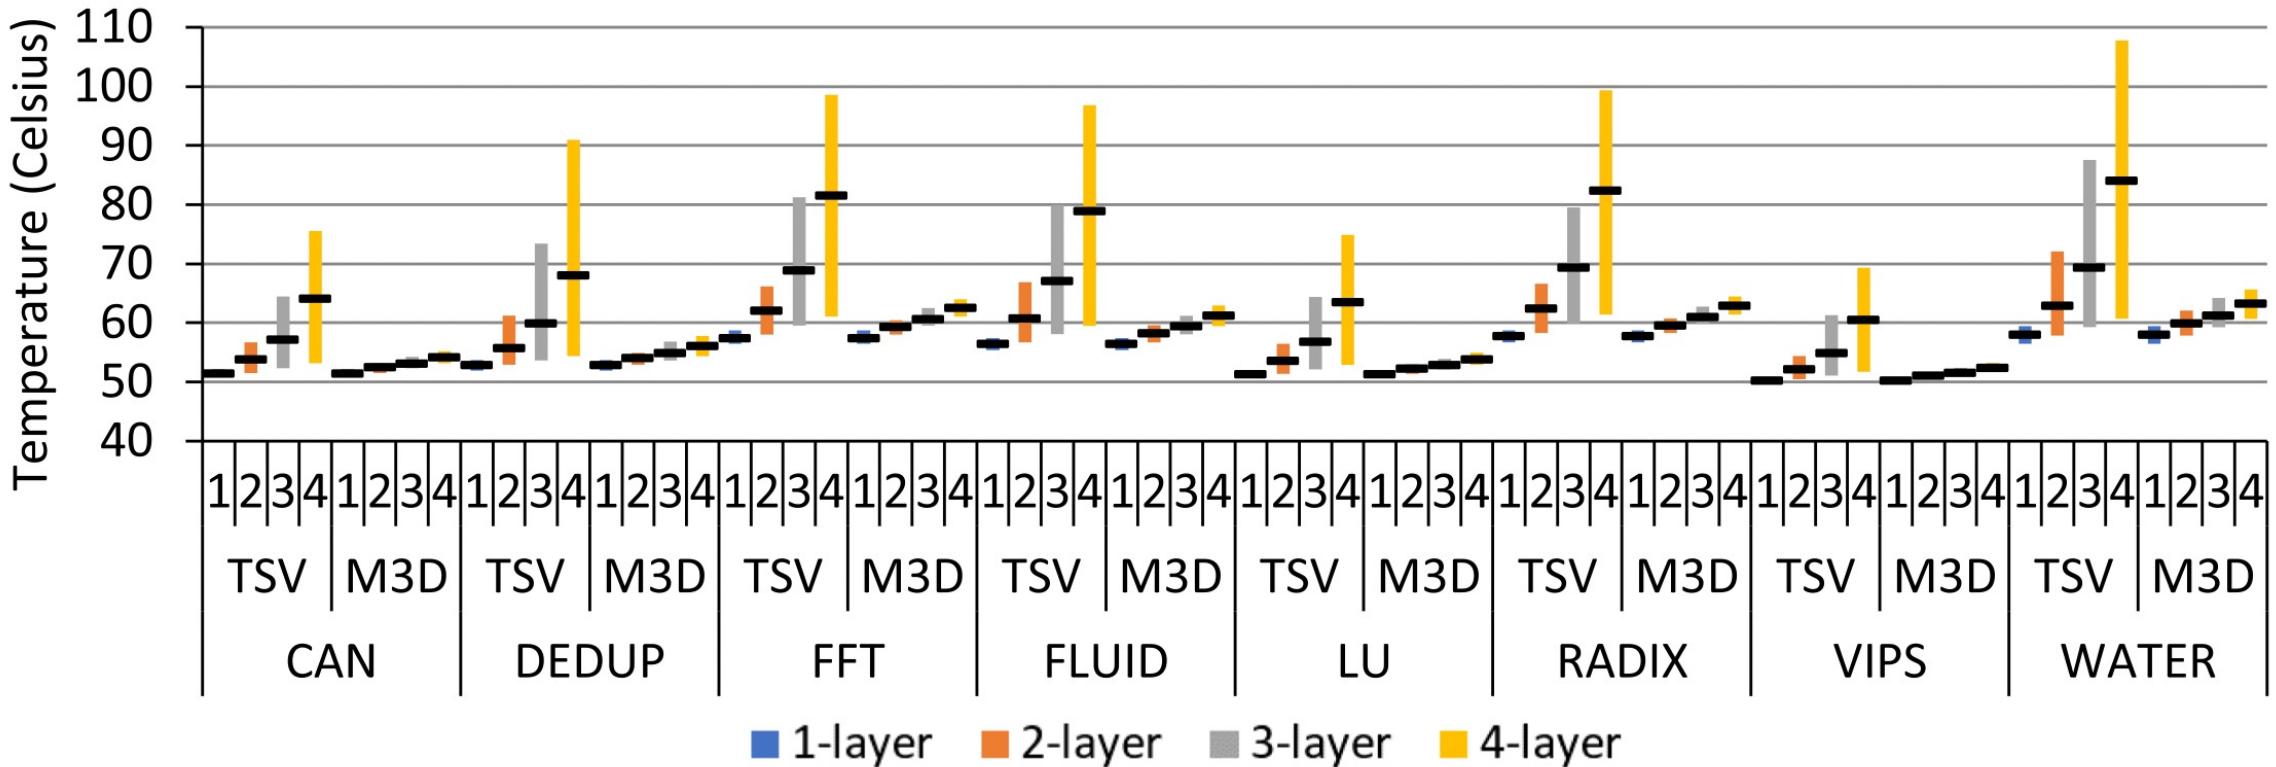


# Fully manycore design using M3D



- Routers and cores can extend over 2 tiers
- Faster logic blocks
- Better performance

# Temperature in 3D



- Layer farthest from the sink has highest temperature
- M3D has better thermal profile than TSV

# How to get power & temperature (2)?

- Cycle-accurate simulators for power
  - CPU: Gem5 + McPat
  - GPU: GPGPU-Sim + GPUWattch
- Allows us to run C/Cuda code
- Simulators for temperature
  - HotSpot, 3D-ICE



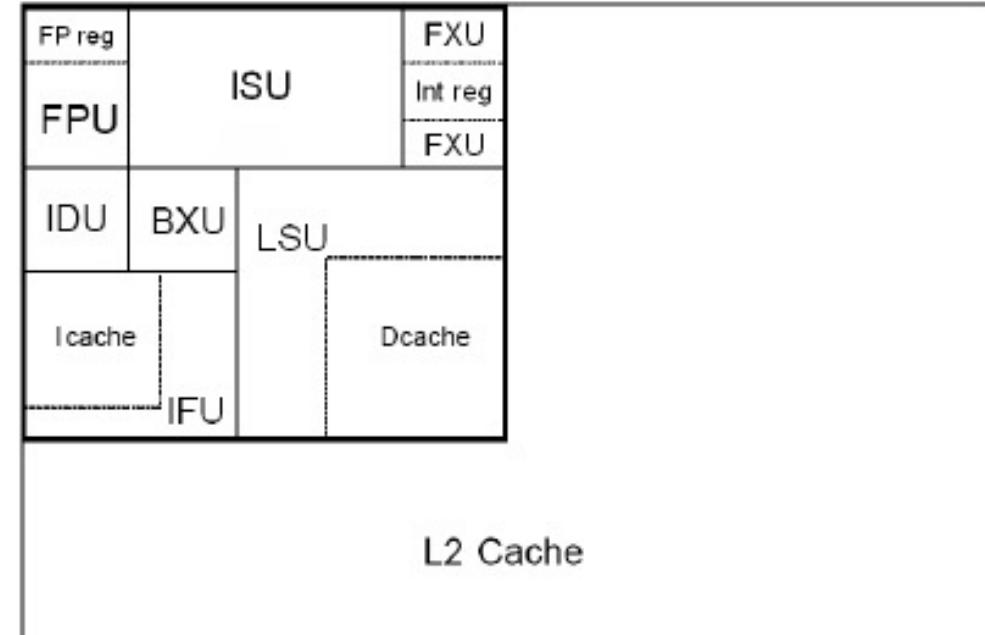
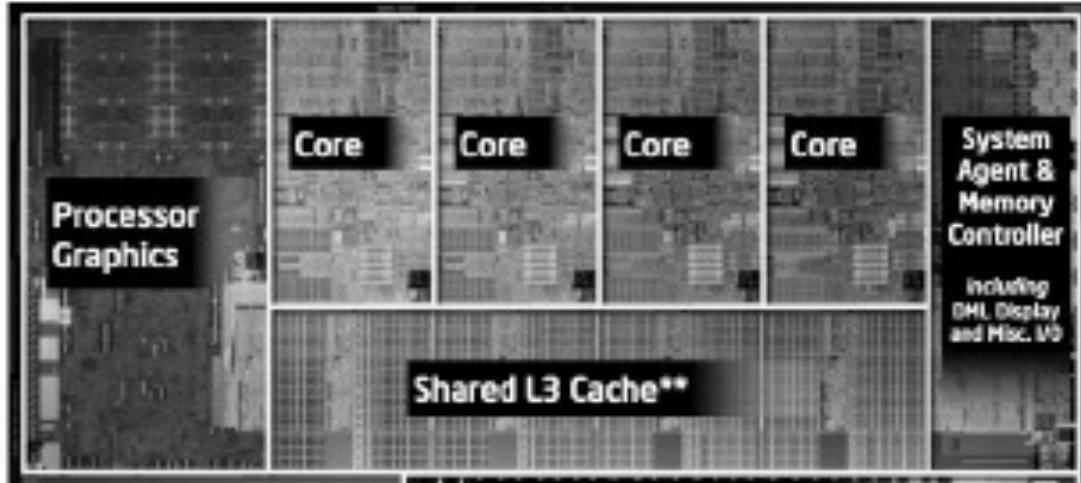
<https://www.gem5.org>

# 3D-ICE thermal simulator

- <https://github.com/esl-epfl/3d-ice>
- Build custom floorplan (.flp)
- Provide technology file (.stk)
- Run 3D-ICE assuming chip dimension of 10mm x 10mm
- Get temperature for various configurations

```
Core0 :  
  
    position      0,      0 ;  
    dimension    2500, 3500 ;  
  
    power values 12.5, 14.0, 10.5, 13.5, 12.5 ;  
  
Core1 :  
  
    position      2500,     0 ;  
    dimension    2500, 3500 ;  
  
    power values 12.5, 14.0, 10.5, 13.5, 12.5 ;  
  
Core2 :  
  
    position      5000,     0 ;  
    dimension    2500, 3500 ;  
  
    power values 12.5, 14.0, 10.5, 13.5, 12.5 ;
```

# Floorplanning



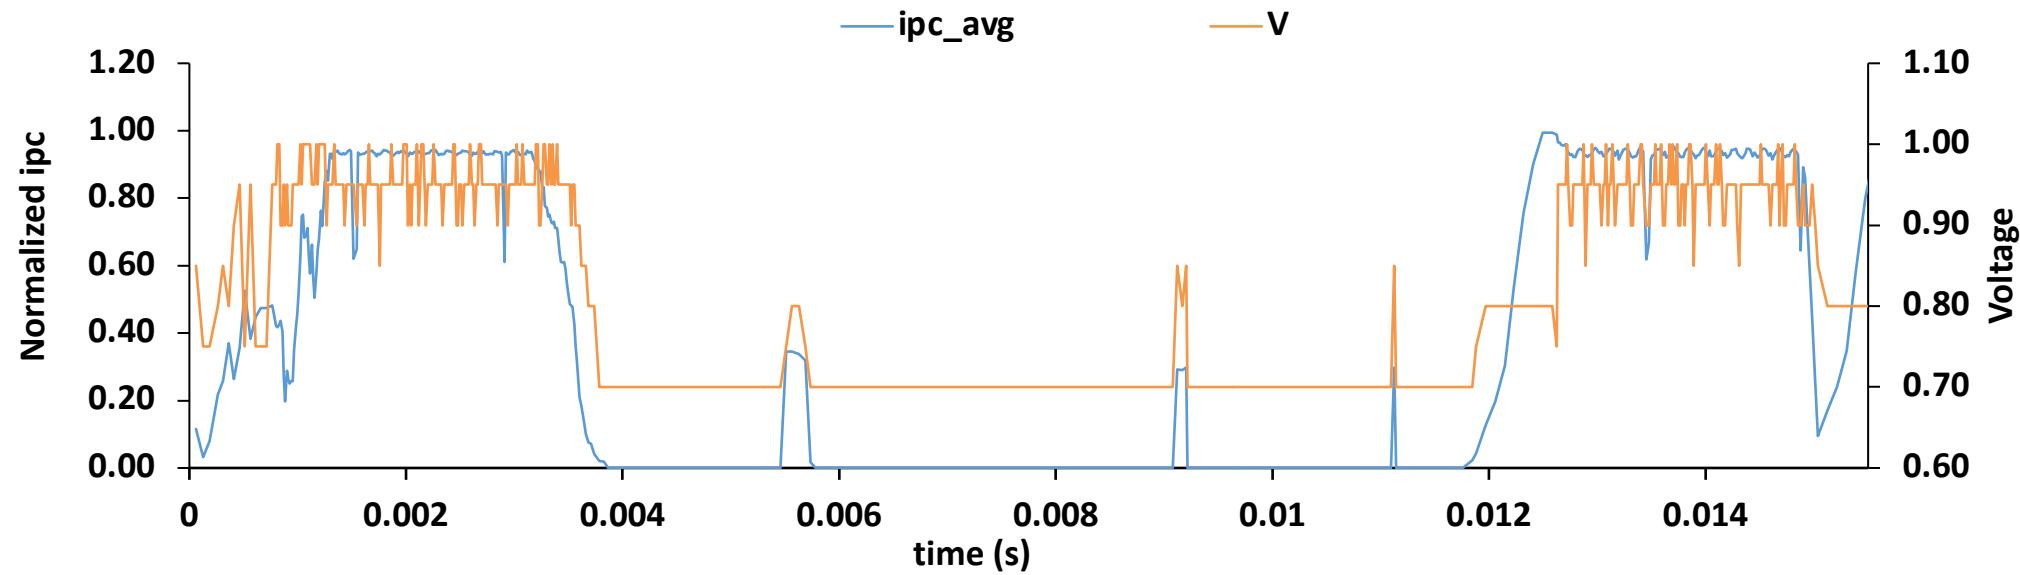
- **Physical placement of cores/logic blocks**
  - Necessary for thermal simulations
  - HotSpot, 3D-ICE

# Power management recap

- Power/Energy is important, especially for embedded systems
  - Mobile and portable devices: Smartphones, Laptops
  - Other devices: Wearables, Medical devices, etc.
- Lower power usage = Better battery life
- More power = More heat

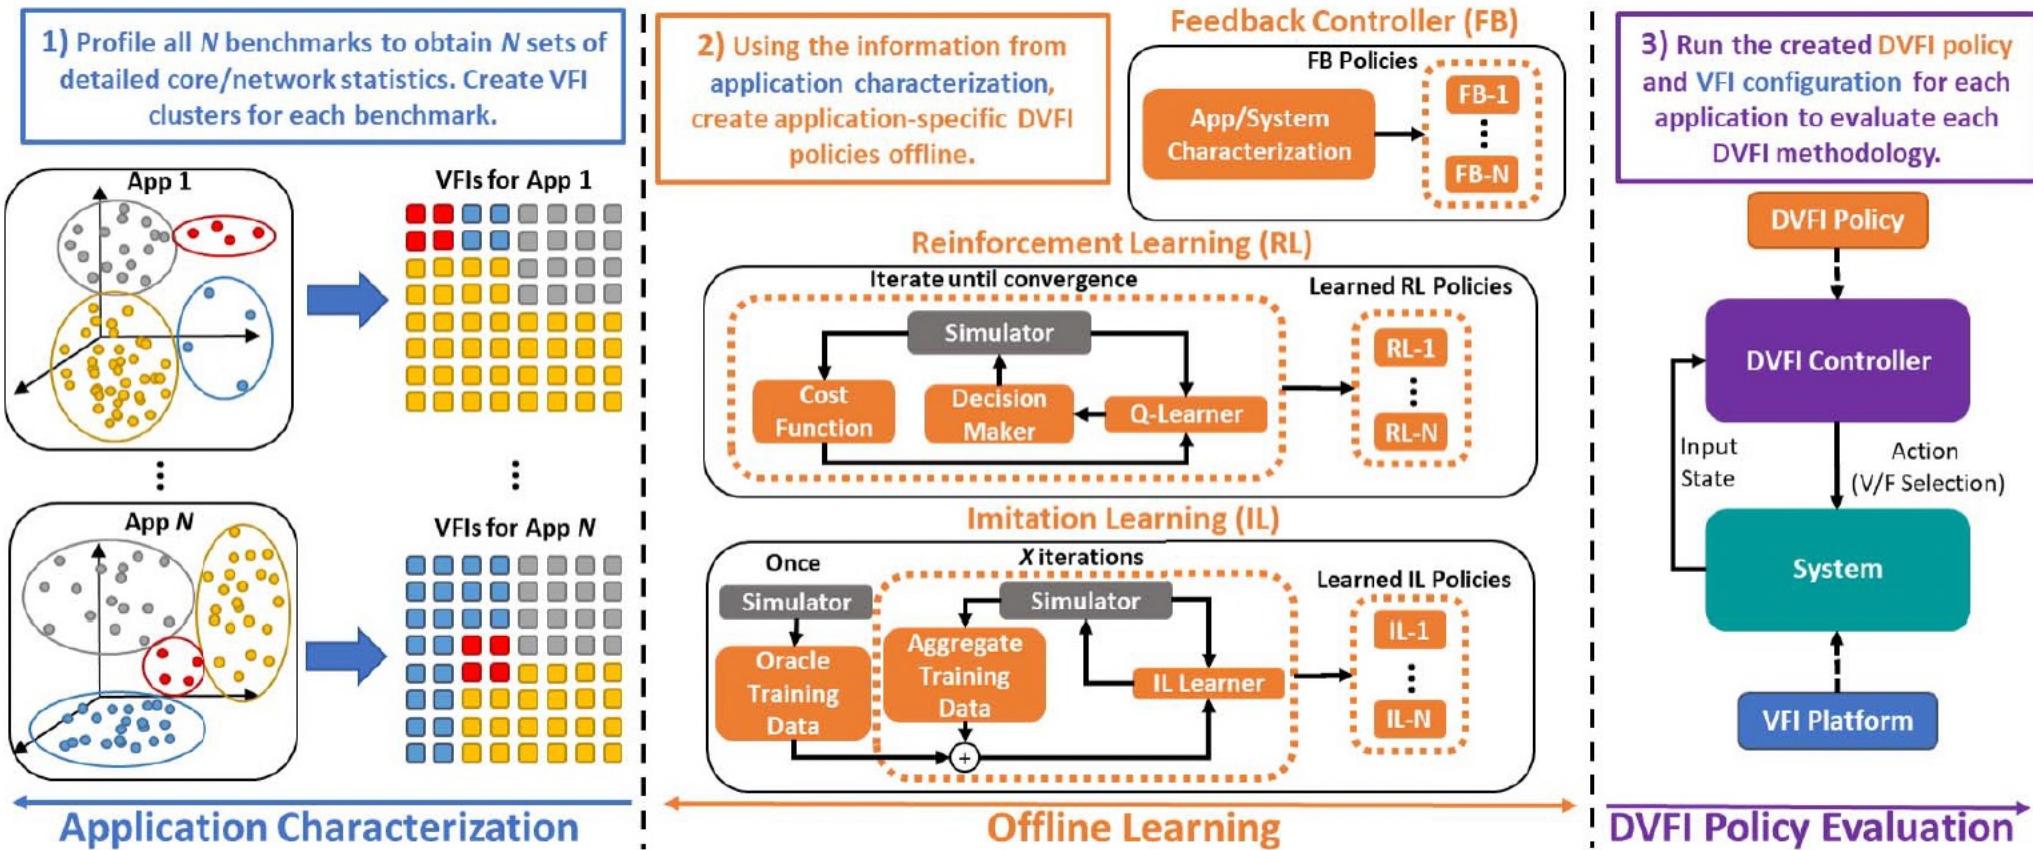


# DVFS recap



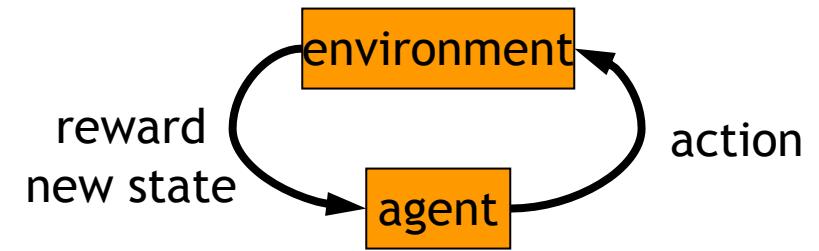
- IPC for FFT with respect to time
- Voltage closely matches the behavior
- ML policy used for prediction

# Overall methodology

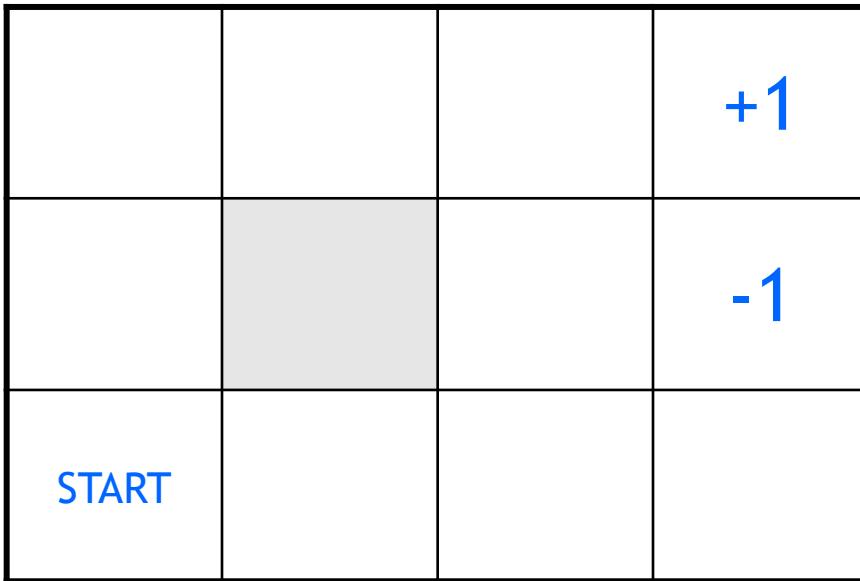


# Reinforcement learning

- Supervised learning
  - classification, regression
- Unsupervised learning
  - clustering
- Reinforcement learning
  - more general than supervised/unsupervised learning
  - learn from interaction w/ environment to achieve a goal



# RL example

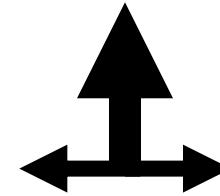


actions: UP, DOWN, LEFT, RIGHT

UP

80%  
10%  
10%

move UP  
move LEFT  
move RIGHT

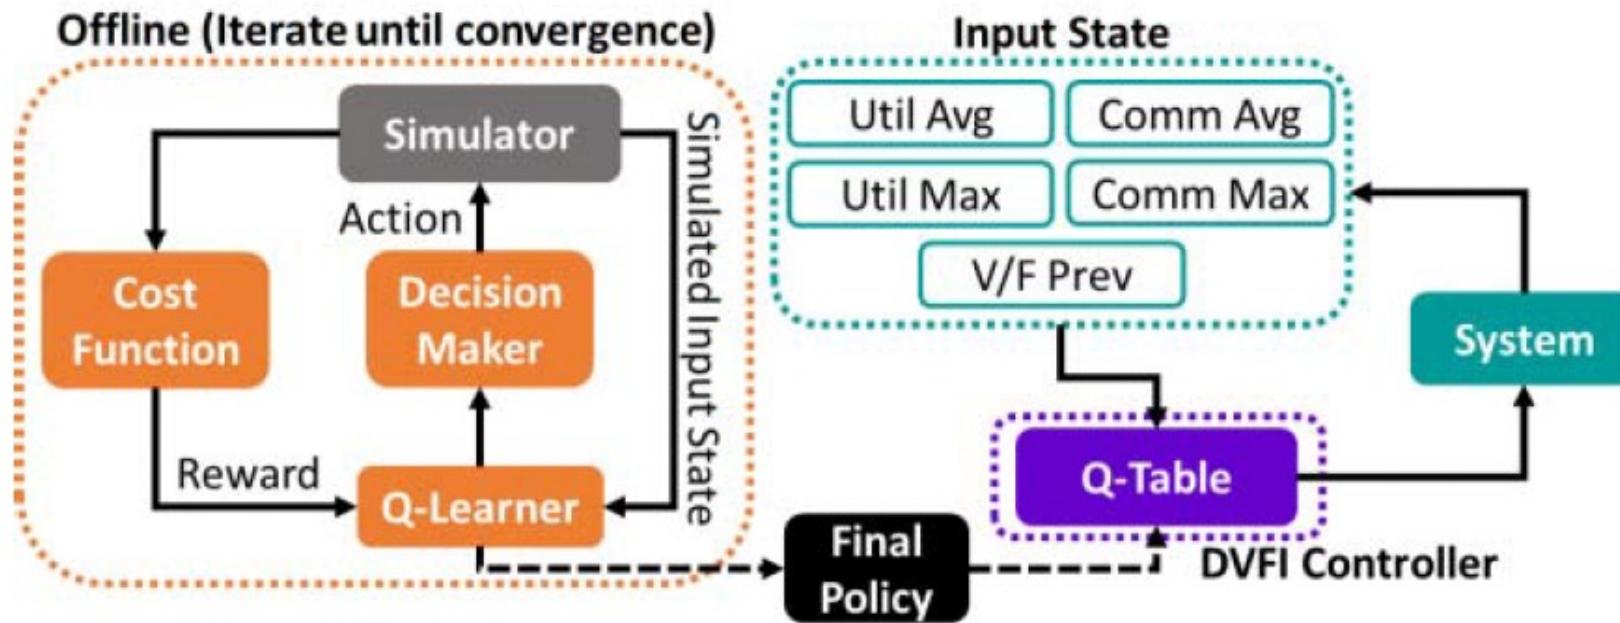


- Maze solving problems
- reward +1 at [4,3], -1 at [4,2]
- reward -0.04 for each step
- what's the strategy to achieve max reward?

<https://www.youtube.com/watch?v=pc-H4vyg2L4>

[https://deeplizard.com/learn/playlist/PLZbbT5o\\_s2xoWNVdDudn51XM8lOuZ\\_Njv](https://deeplizard.com/learn/playlist/PLZbbT5o_s2xoWNVdDudn51XM8lOuZ_Njv)

# Reinforcement learning for power management



- Which  $\langle V, F \rangle$  pair for a certain time window?
- RL to learn the best decision

# Imitation learning (1)

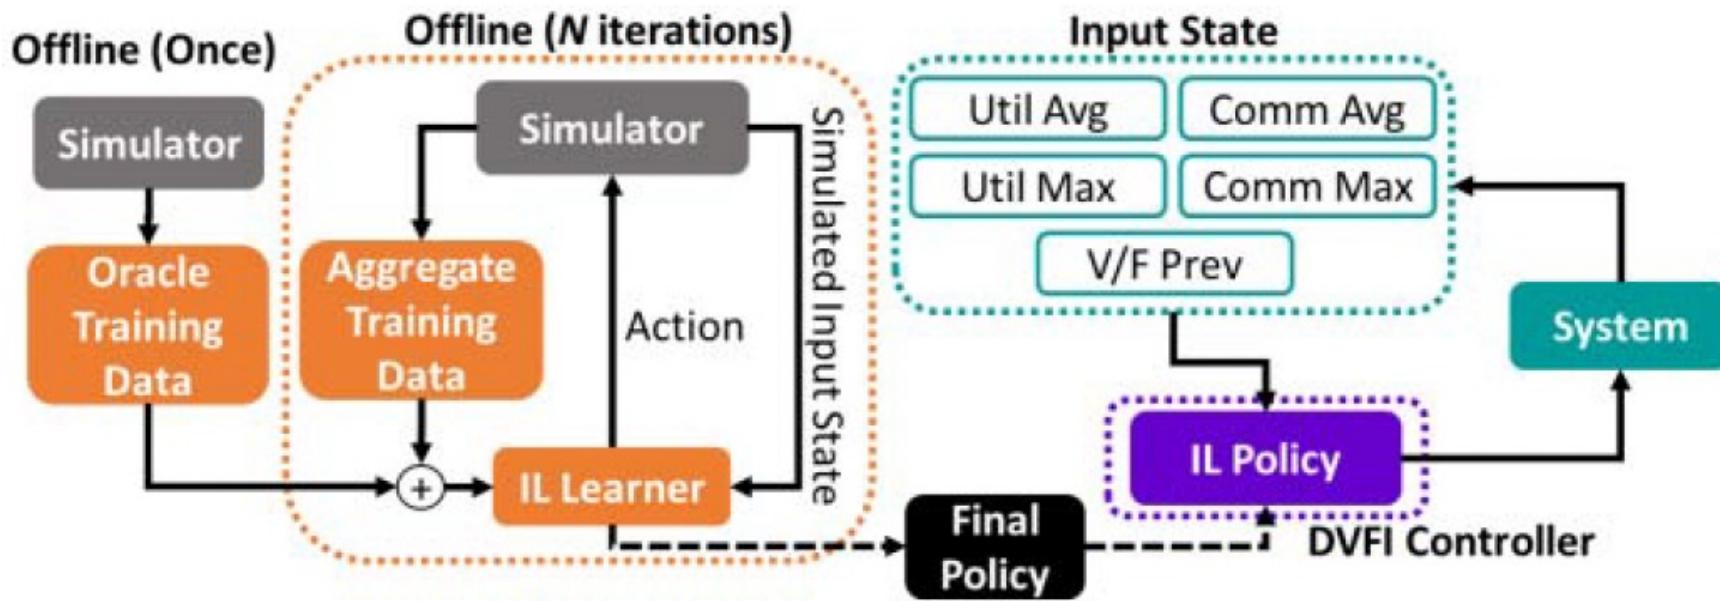


- Imitate the expert
- Simpler and faster than RL

# Imitation learning vs Reinforcement learning

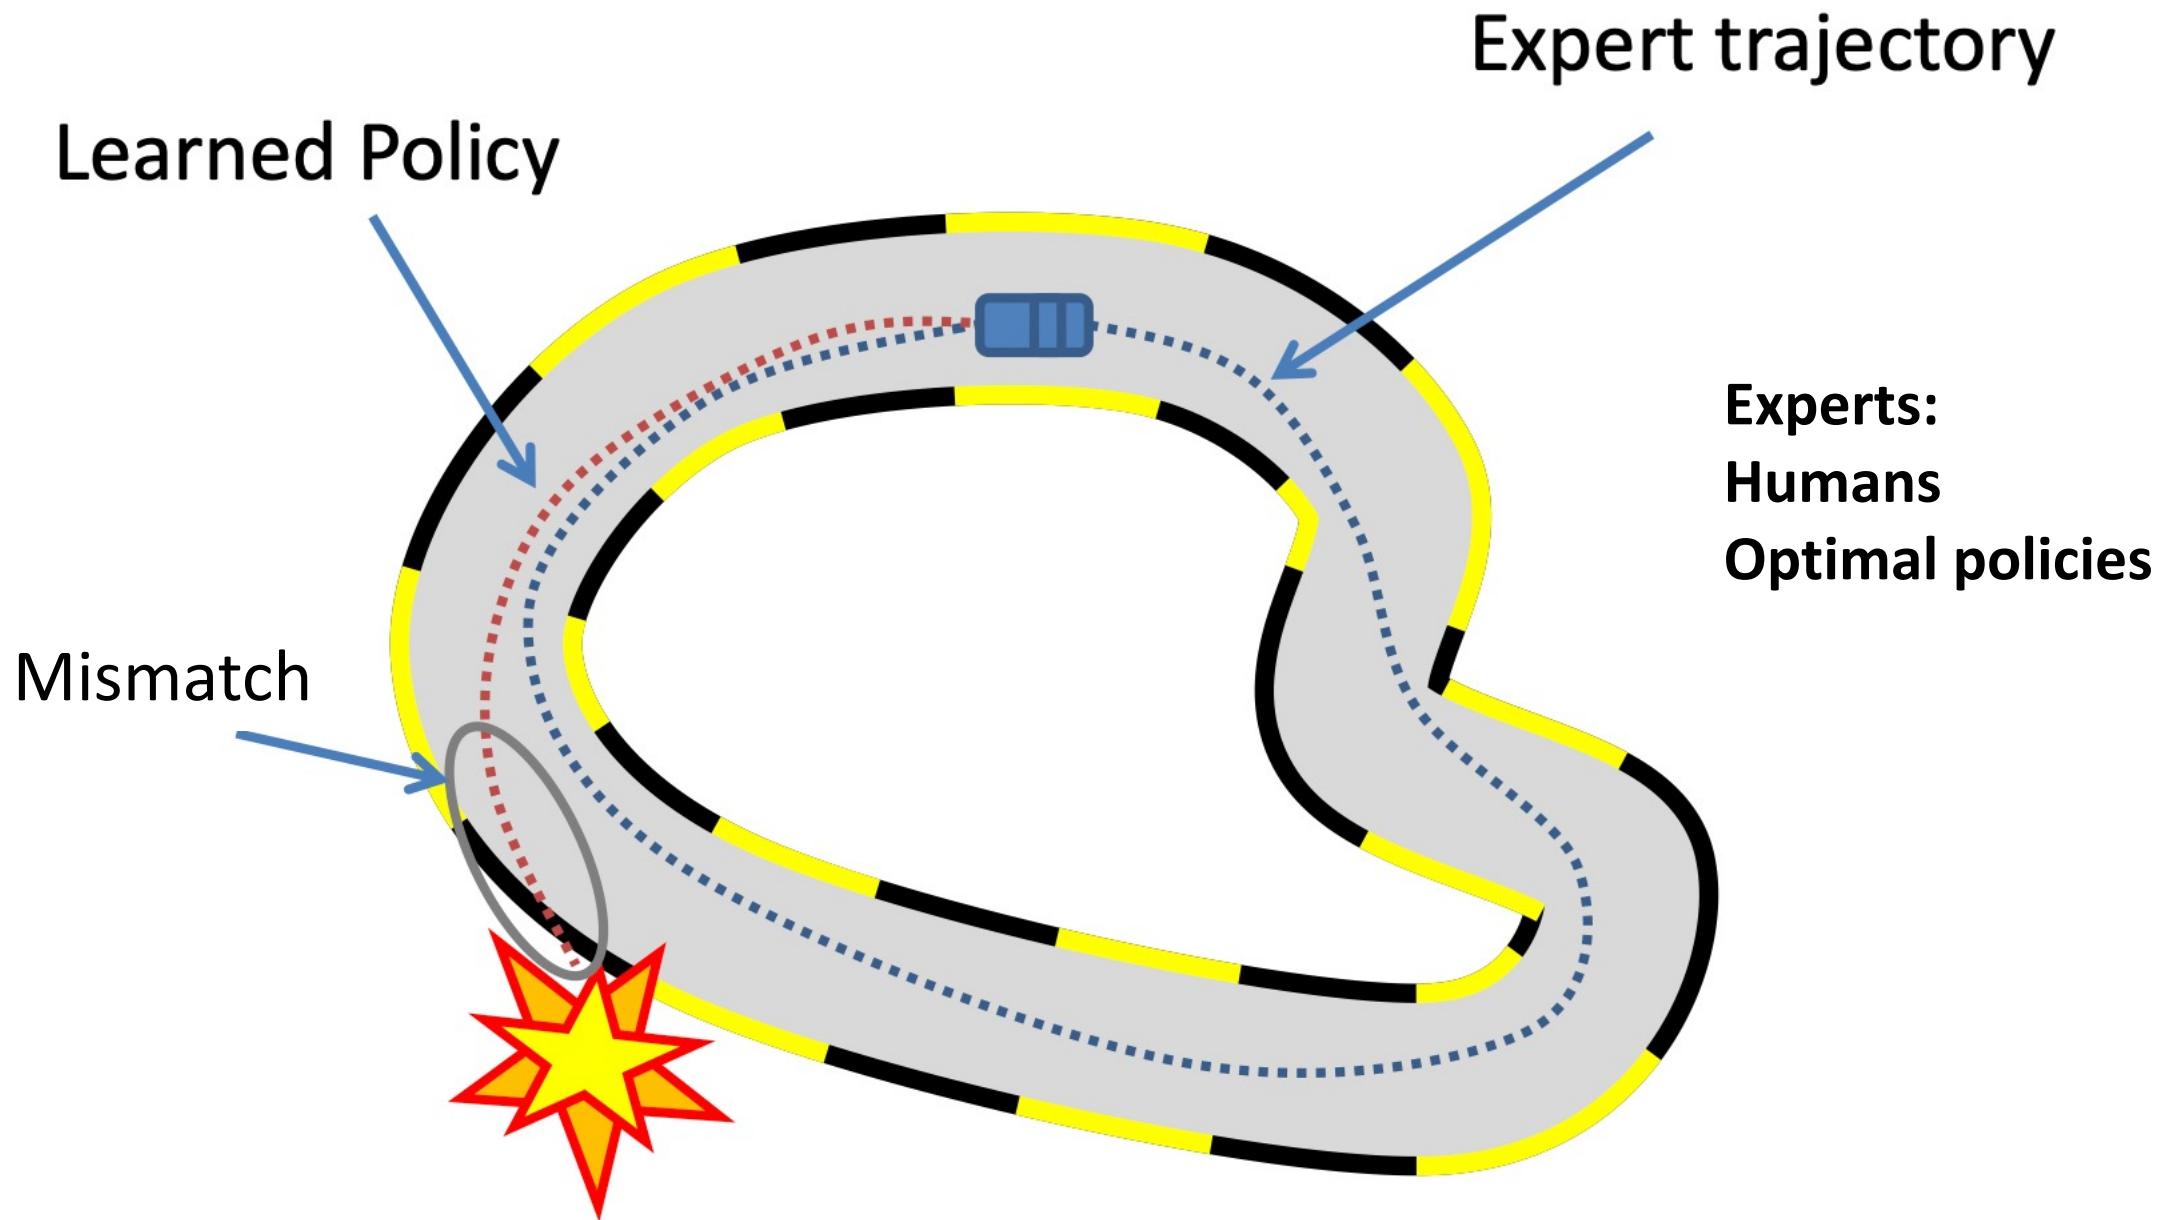
- **Reinforcement Learning:** Learning policies guided by sparse rewards, e.g., win the game.
  - Good: simple, cheap form of supervision
  - Bad: High sample complexity
- In simulation, where we can afford a lot of trials, easy to parallelize
- Not in robotic systems:
  - Action/reward takes long
  - we cannot afford to fail (e.g., Offroad navigation!)
  - safety concerns
- Imitation learning is useful when it is easier for the expert to demonstrate (or oracle creation can be done easily)

# Imitation learning for VFI

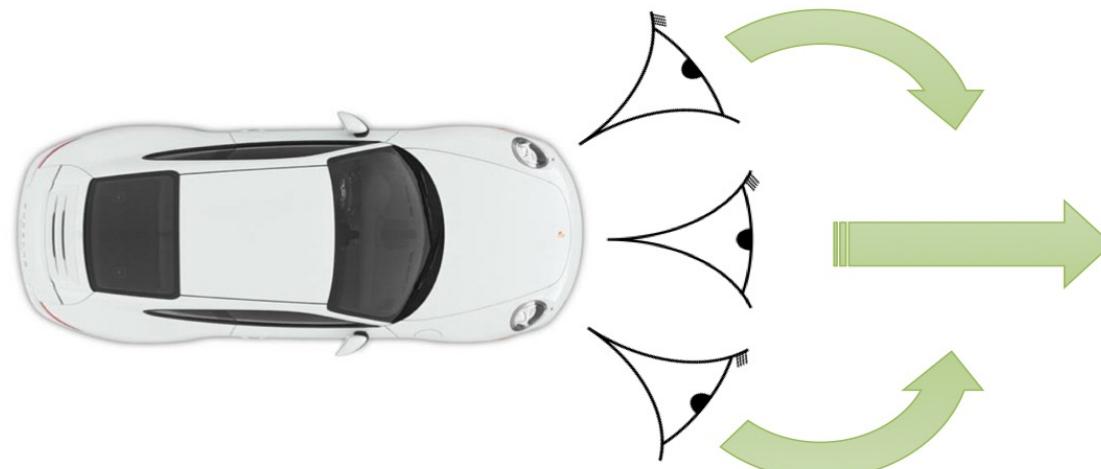
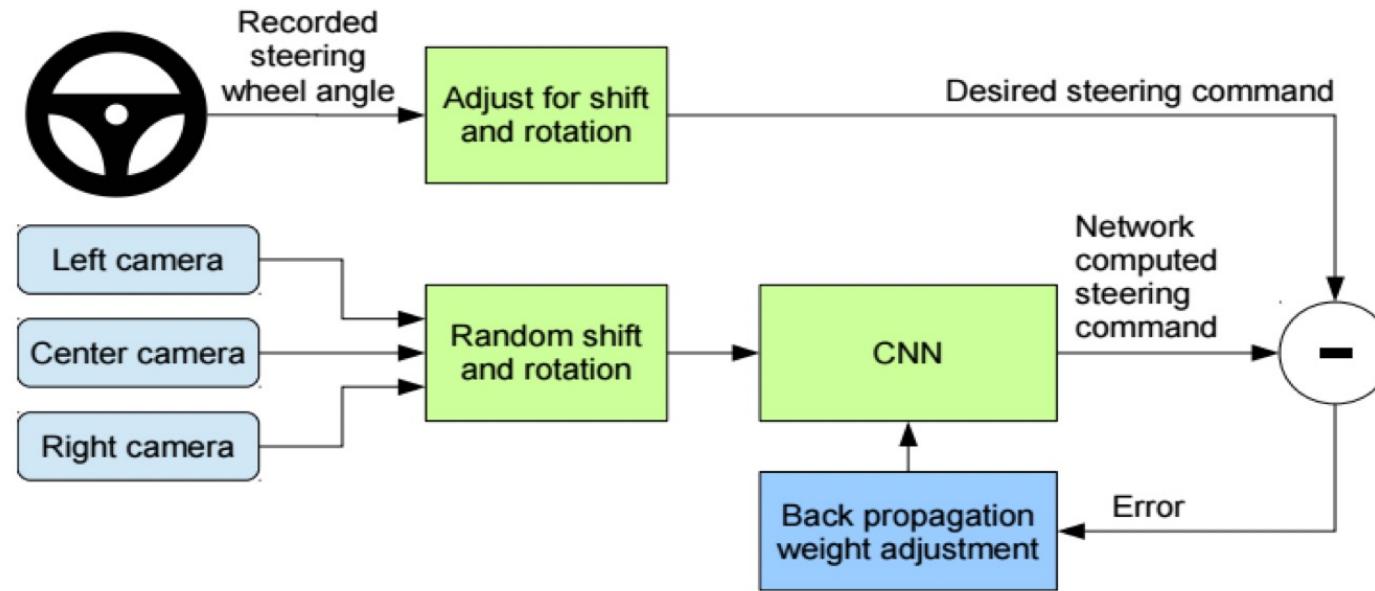


- Which  $\langle V, F \rangle$  pair for a certain time window?
- IL to learn the best decision
- Where is the expert?

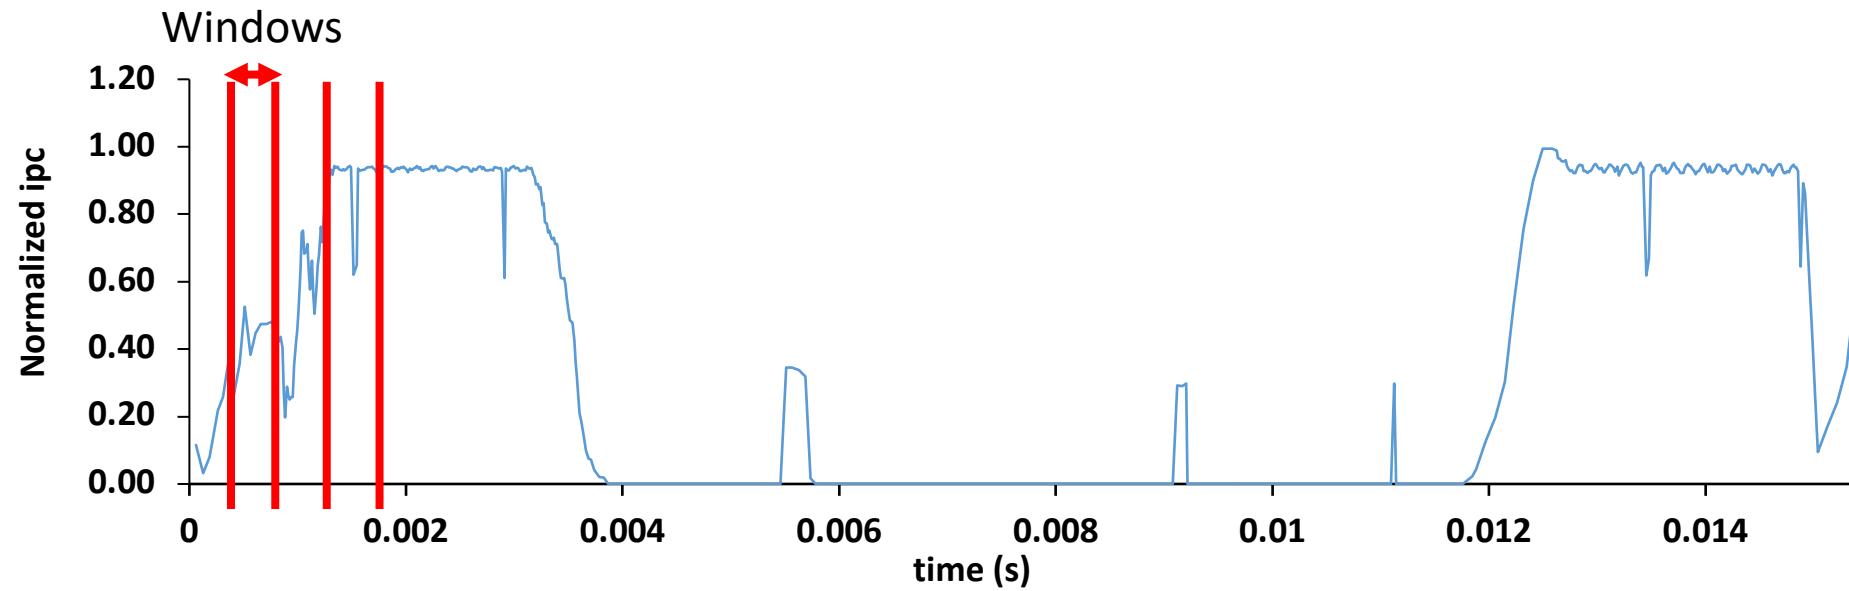
## Imitation learning (2)



# Imitation learning (3)

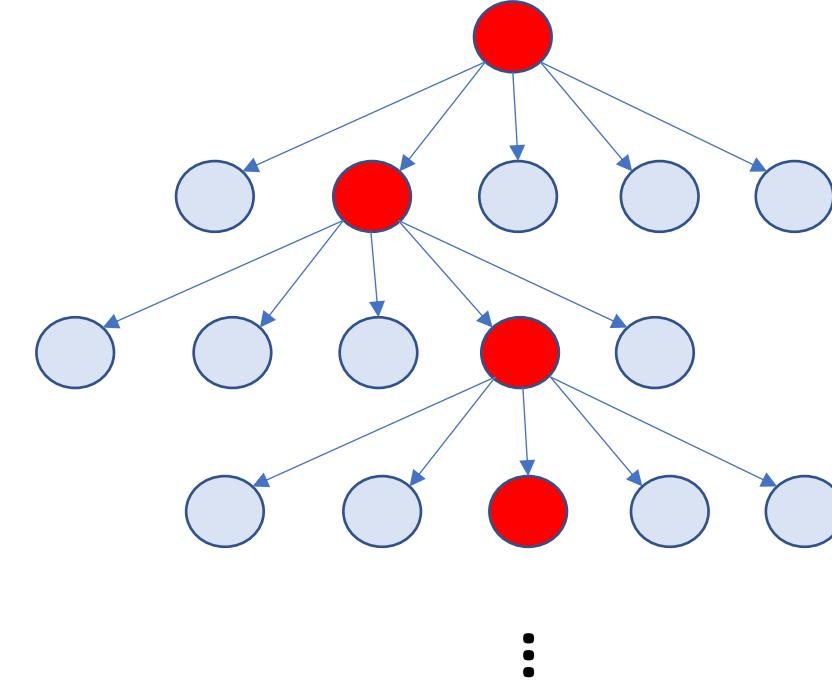
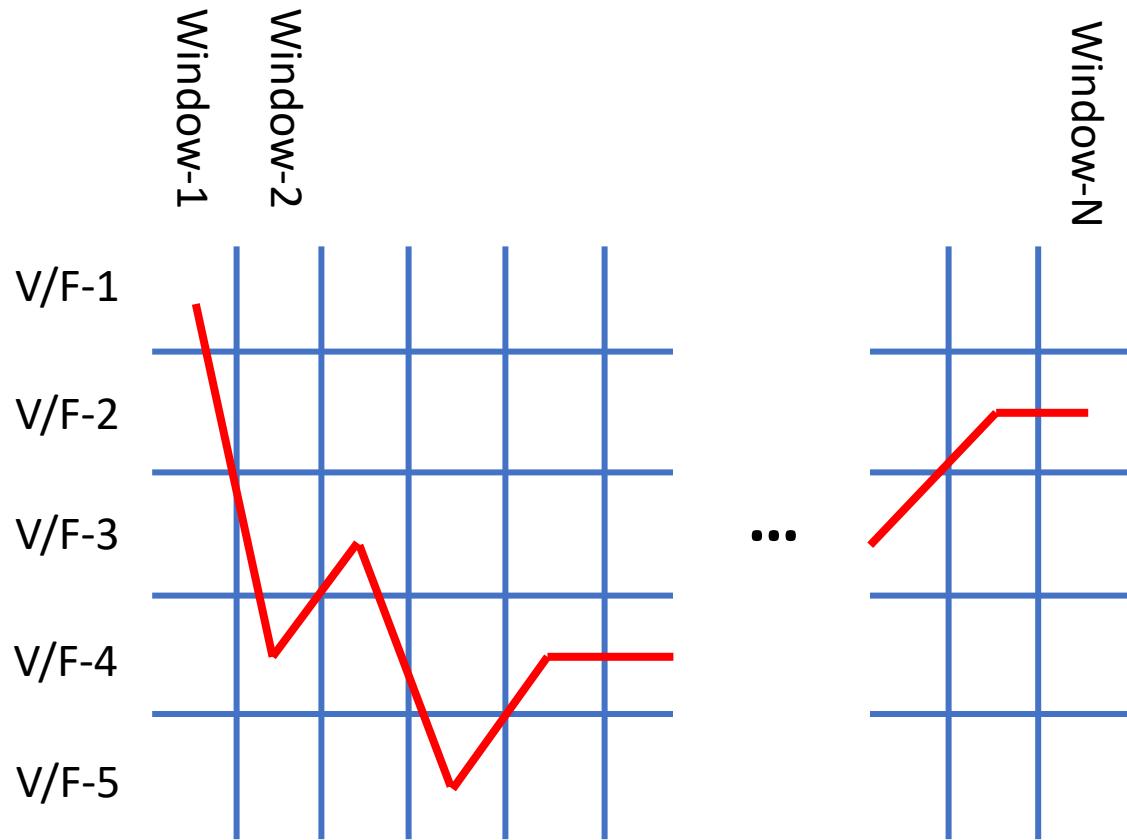


# Creating the oracle for IL-based DVFI(1)



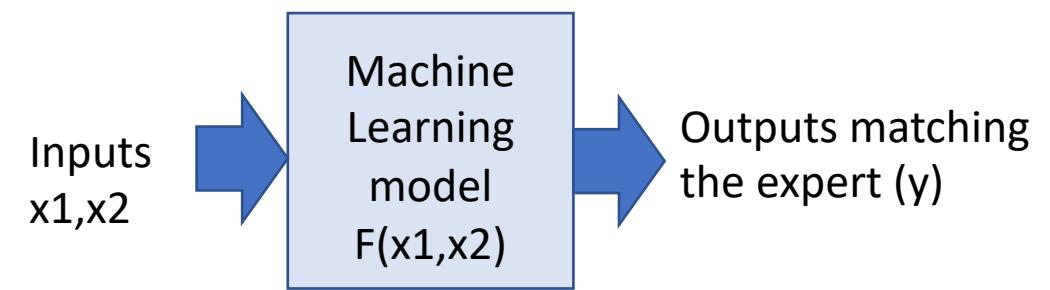
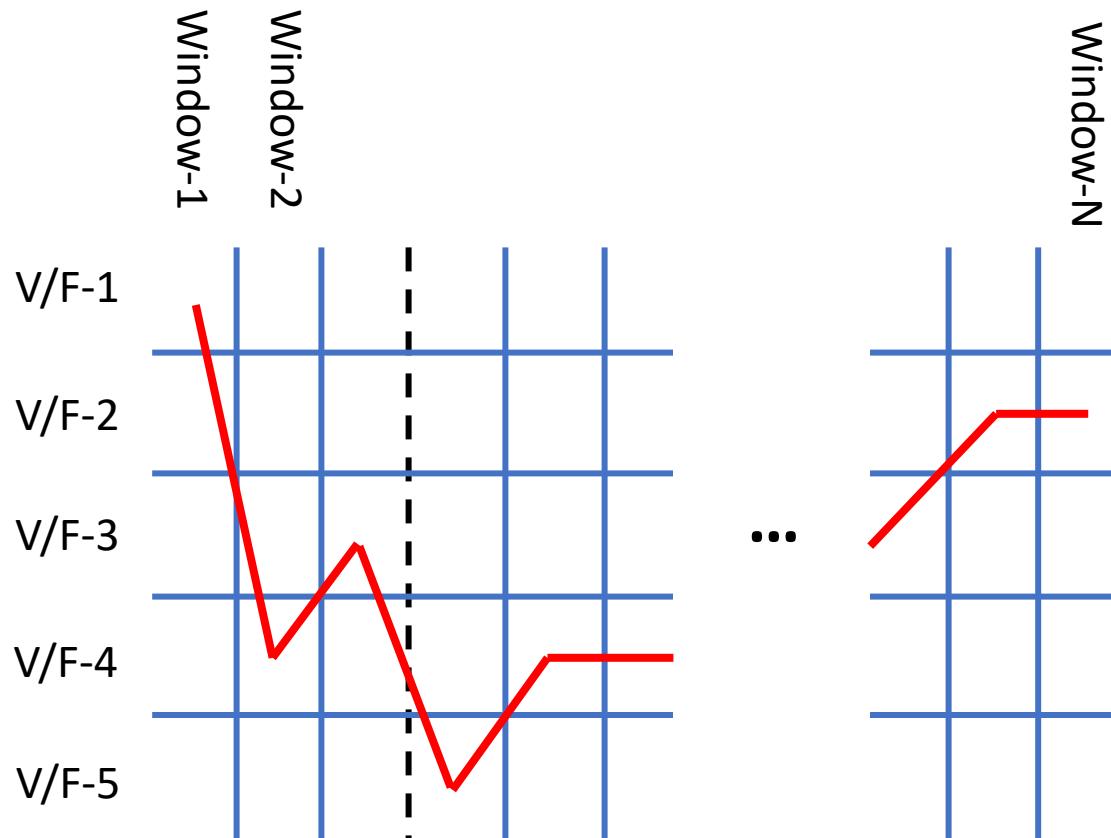
- Window-wise data for different V/F using Gem5
- Impossible to create oracle considering all windows
- Create the oracle using smaller segments

# Creating the oracle for IL-based DVFI(2)

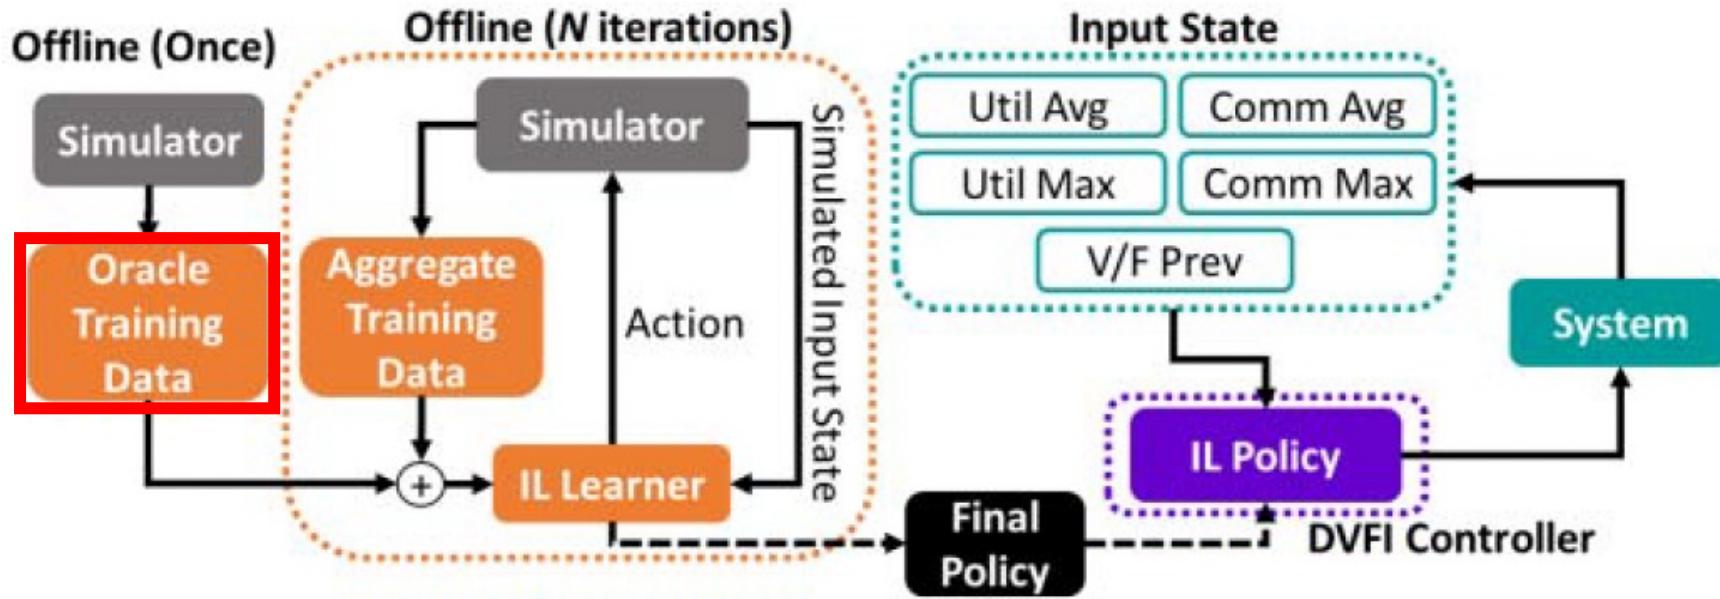


- Find the right set of V/F for each window to reduce overall power

# Imitating the Expert

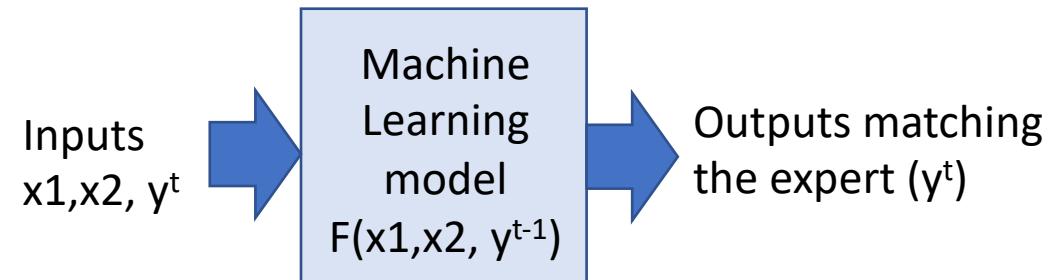
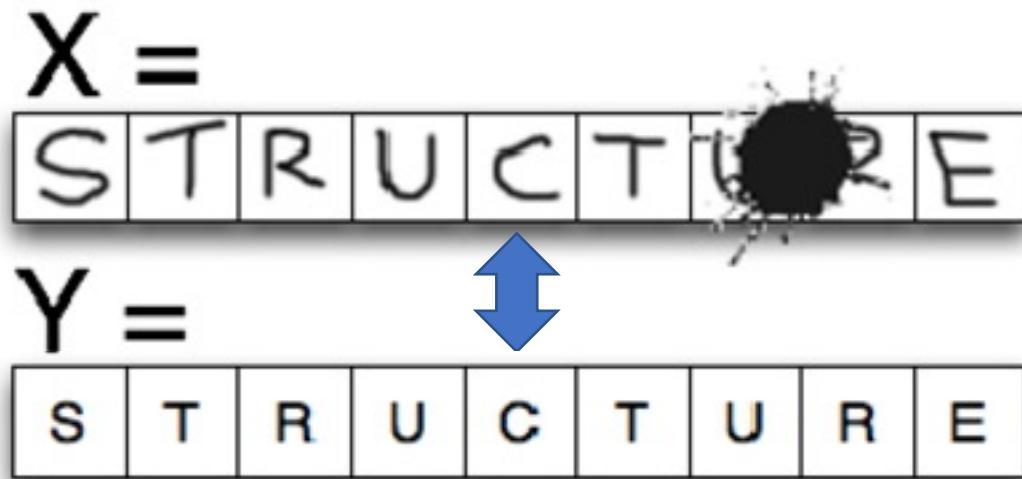


# IL for calculating V/F



- We have the oracle now
- Use traditional ML to learn using oracle as ground truth
- IL learner can be linear regression or any other model

# Structured Prediction

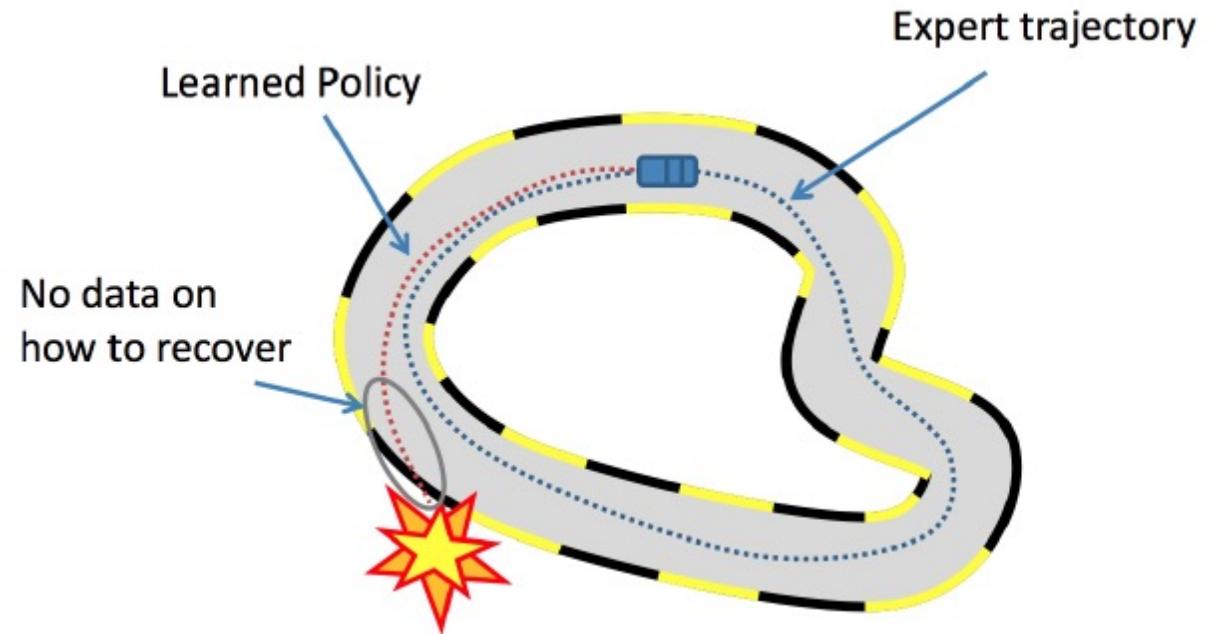


- Does V/F of previous window affect V/F of current window?

# Dagger algorithm(1)

```
Initialize  $\mathcal{D} \leftarrow \emptyset$ .  
Initialize  $\hat{\pi}_1$  to any policy in  $\Pi$ .  
for  $i = 1$  to  $N$  do  
    Let  $\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}_i$ .  
    Sample  $T$ -step trajectories using  $\pi_i$ .  
    Get dataset  $\mathcal{D}_i = \{(s, \pi^*(s))\}$  of visited states by  $\pi_i$   
        and actions given by expert.  
    Aggregate datasets:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$ .  
    Train classifier  $\hat{\pi}_{i+1}$  on  $\mathcal{D}$ .  
end for  
Return best  $\hat{\pi}_i$  on validation.
```

**Algorithm 3.1:** DAGGER Algorithm.



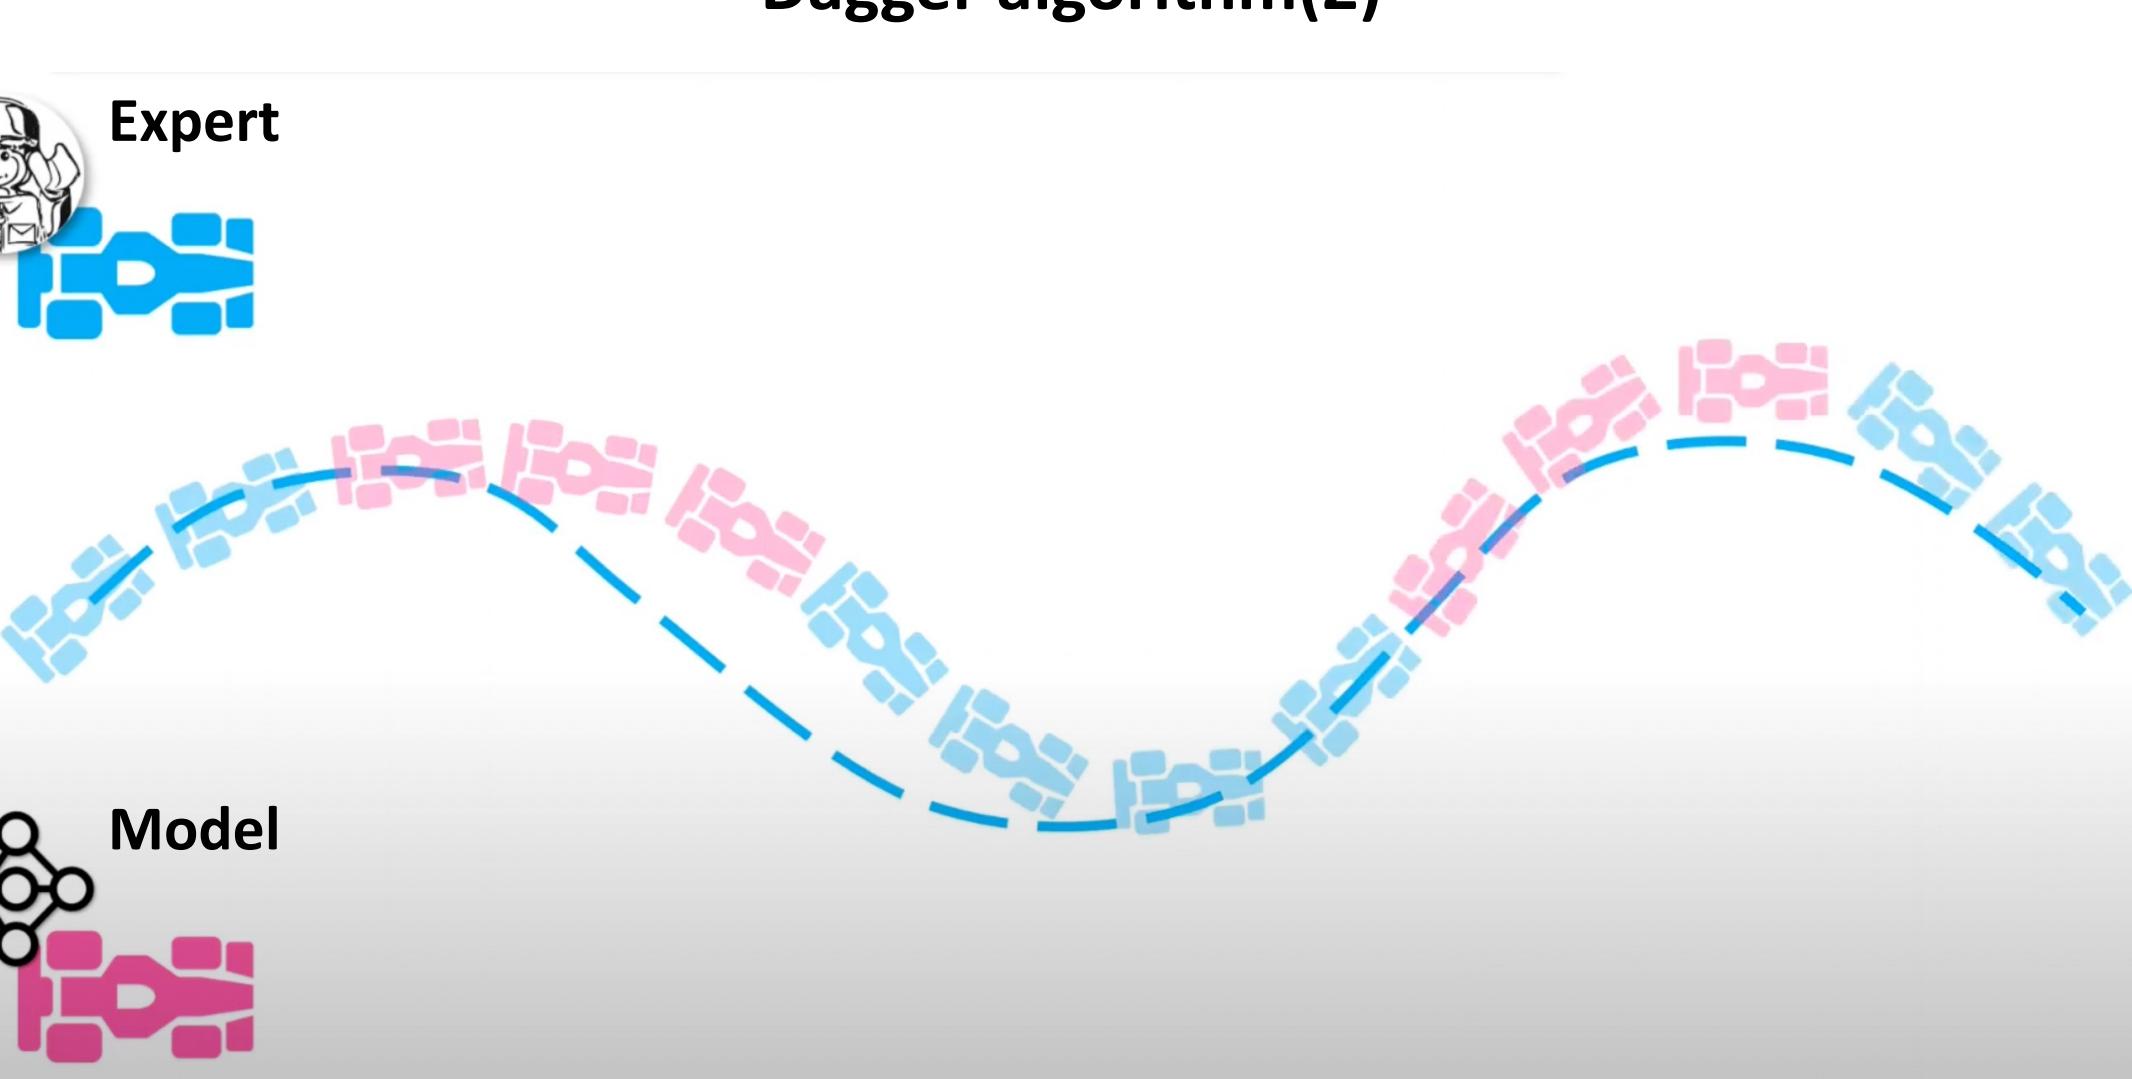
## Dagger algorithm(2)



Expert



Model



- Sometimes let the model make mistakes
- Correct them using the expert